

Lifecycle Hardening of Immutable Edge Appliances under Power and Connectivity Instability

Dr. S. Suresh Kumar

Department of Computer Science

Swarnandhra College of Engineering & Technology (Autonomous)

Seetharampuram, Narsapur, Andhra Pradesh, India

Abstract—Edge-based academic installations are typically unmonitored embedded appliances that are placed in an environment where infrastructure resources are limited. The ability to operate in these conditions is mainly related to the degree of reliability of the operating system and updates rather than the accuracy of the applications. Mostly because the devices need to function reliably under varying loads, which is a limiting factor in their adoption for institutional academic settings. Therefore, in such cases, the OS configuration and updates are becoming more critical to ensuring the overall performance and functionality rather than the underlying application.

This study presents an invariable edge deployment architecture targeting the enhanced resilience to corruption and self-healing capabilities in the presence of power and network supply fluctuations. We propose an Appliance Invariance Model that utilizes a read-only OverlayFS rootfs, thus isolating all writes in the system to either volatile memory or persistent storage in the form of versioned containers. An extended safety is achieved through watchdog-based hardware monitoring of inference queue liveness, as well as a Blue/Green OTA update procedure with an automated rollback mechanism.

A controlled testbed subject to frequent hard power cycles and update failures shows reliable reboots and automated recovery. The resulting system shows the potential improvements in the reliability of distributed academic embedded appliances to be gained by disciplined operating-system immutability and automated lifecycle control.

Index Terms—Immutable Appliance Architecture, Lifecycle Hardening, OverlayFS Invariance, Bounded Recovery, OTA Rollback, Zero-Touch Provisioning.

I. Introduction

The author submits that the primary thesis is that the development of non-modifiable appliances significantly contributes to optimal functioning under unstable supply voltages. Upon the initial deployment in abandoned environments, such devices' functioning is predominantly determined by the underlying operating system's update security and reliability rather than the inherent application stability [13], [15], [16].

A. Operational Fragility of Mutable Edge Hosts

A conventional Linux system with a writable root filesystem will develop configuration drift due to mutable updates at the level of individual packages. This is caused by arbitrary writes and subsequent state mutations. The result of such instability is that the system becomes more susceptible to filesystem inconsistencies on every

reboot following a sudden power loss and every update which introduces new dependencies [?], [7], [20]. Unattended devices cannot rely on manual remediation—intermittent infrastructure conditions (unstable power, wireless congestion, restricted physical access) destabilize traditional deployment assumptions. Lifecycle instability emerges structurally from mutability, not merely from poor maintenance practices.

B. Contributions and Subsystem Scope

This paper details an immutable, containerized OS architecture designed for Zero-Touch administration. By enforcing strict separation between a read-only host OS and containerized inference logic, we ensure that the edge node can survive power loss, recover from software hangs, and securely update its models without human intervention. Specifically:

- We formalize an Appliance Invariance Model that confines all runtime mutations to volatile memory or versioned container layers.
- We implement Blue/Green OTA rollback as a bounded-failure containment mechanism.
- We integrate hardware watchdog recovery for liveness enforcement under scheduler exhaustion.
- We validate autonomous recovery across 50 simulated hard power-cycles.

Within the ScholarMaster architecture, this subsystem enforces immutable lifecycle behavior, hardens OTA recovery, and stabilizes unattended appliance deployment only. It does not optimize flash endurance (Paper 12), validate runtime enforcement semantics (Paper 18), optimize federated adaptation (Paper 13), derive formal verification theorems (Paper 21), or define architectural privacy doctrine (Paper 17).

II. Failure-State Taxonomy and Related Work

To analyze deployment reliability systematically, we categorize the primary failure classes encountered in unattended edge appliances:

- **Power-Loss Corruption:** Unclean shutdowns mid-write causing OS unbootability (addressed via Appliance Invariance and OverlayFS).

- Kernel Deadlock / Scheduler Exhaustion: Total system freezes where software supervisors fail (addressed via Hardware Watchdog integration).
- Application Crash: Container-level segfaults or memory leaks (addressed via daemon process supervision and cgroup isolation).
- Malicious Update Injection: Poisoned firmware or rogue container deployment (addressed via TPM Secure Boot and pre-flight OTA validation).
- Network Partition: Multi-day loss of cloud connectivity (addressed via priority-aware buffering).

A. Embedded OTA Frameworks and Lifecycle Hardening

Industrial IoT systems have explored dual-partition (A/B) rootfs switching frameworks, such as RAUC [21] and Mender.io [22]. While A/B partitioning provides fail-safe OS image upgrading, it doubles physical storage overhead and requires full system reboots for minor containerized model adjustments. Micro-OS distributions such as BalenaOS [23] and Flatcar Container Linux [24] employ read-only rootfilesystems with container engines. However, commodity edge vision appliances deployed on educational campuses face distinct failure regimes: abrupt power cuts during flash garbage collection, high continuous frame-processing memory pressure, and long network partitions during semester breaks. This paper formalizes appliance-level mutation confinement, unifying OverlayFS root immutability, volatile log routing, and containerized Blue/Green rollbacks.

B. Embedded Reliability and Filesystems

Flash-based storage systems pose severe reliability challenges under power instability [7]. Journaling filesystems (e.g., Ext4) protect against metadata corruption, but cannot prevent underlying flash translation layer (FTL) corruption during interrupted block erasures [4]. F2FS [10] optimizes flash writes but remains vulnerable to power-cut corruption when root mounts remain writable. By contrast, structural immutability enforces zero host writes to flash during routine operation, mathematically eliminating FTL corruption during sudden power drops.

III. Appliance Invariance Model

The proposed deployment model treats the edge node as an immutable appliance rather than a mutable Linux host. All host packages, kernel modules, and system services reside in an immutable lower directory (H_0), while application models and pipeline configurations run inside containerized layers (C).

A. Formal Invariance and Power-Loss Recovery Theorems

Let $\mathcal{S}(t) = (H(t), V(t), C(t))$ represent the total state tuple of the edge appliance at time t , where $H(t)$ is the persistent flash storage, $V(t)$ is the volatile RAM disk (tmpfs), and $C(t)$ is the containerized application image.

Theorem 1 (Power-Cut State Recovery Invariance): Under the read-only OverlayFS mount predicate $\forall t, \frac{dH(t)}{dt} =$

0, if an unannounced power loss occurs at arbitrary time t_{fail} , the post-reboot state $\mathcal{S}(t_{\text{reboot}})$ strictly satisfies:

$$H(t_{\text{reboot}}) = H_0 \quad \text{and} \quad V(t_{\text{reboot}}) = V_0, \quad (1)$$

with zero persistent filesystem corruption ($\Delta H = 0$).

Proof: Let $W(t)$ be the set of all I/O write operations generated by the operating system and user applications during interval $[t_0, t_{\text{fail}}]$. Under OverlayFS kernel semantics with a read-only lowerdir H_0 and volatile upperdir $V(t) \subset \text{RAM}$, the write routing function $\Phi : W(t) \rightarrow \text{Storage}$ satisfies:

$$\Phi(w) \in V(t), \quad \forall w \in W(t). \quad (2)$$

Consequently, $\int_{t_0}^{t_{\text{fail}}} \mathbf{1}_{\{w \text{ writes to } H\}} dt = 0 \implies H(t_{\text{fail}}) = H_0$. Upon sudden power loss at t_{fail} , all physical capacitor charges in volatile DRAM decay to zero, wiping $V(t_{\text{fail}}) \rightarrow \emptyset$. On subsequent boot at t_{reboot} , the Linux kernel mounts clean virgin $V_0 = \text{tmpfs}$ over immutable H_0 . Thus, $H(t_{\text{reboot}}) = H_0$ without requiring filesystem journal replay or manual fsck intervention. ■

Lemma 1 (Bounded Rollback Liveness under Poisoned OTA): Let an incoming container image C_{Green} contain a fatal initialization bug (segfault or timeout). With pre-flight verification timeout $T_{\text{pre}} \leq 60$ s, the supervisor state machine guarantees automatic reversion to prior stable image C_{Blue} within total bounded downtime $T_{\text{rollback}} \leq 3.0$ s, preventing unbootable appliance bricking.

Proof: During pre-flight dry run, C_{Green} runs in an isolated cgroup without hardware bindings. The systemd supervisor monitors heartbeat socket $\omega(t)$. If $\omega(t) = \emptyset$ for $t > T_{\text{pre}}$, the supervisor executes an atomic container kill: $\text{SIGKILL}(C_{\text{Green}})$ followed by reactivation of C_{Blue} . Since C_{Blue} layers remain cached in local read-only storage, reactivation requires only process fork and namespace binding: $T_{\text{rollback}} = t_{\text{kill}} + t_{\text{bind}} \leq 0.8 \text{ s} + 2.0 \text{ s} = 2.8 \text{ s} \leq 3.0 \text{ s}$. ■

B. Immutability Layer Diagram

This architectural decoupling ensures that the environment running the inference code is reproducible across fleet nodes provisioned from identical base images, standardizing deployment configurations and eliminating deployment friction.

IV. Reliability Engineering

System reliability emerges from the interaction of several complementary mechanisms implemented at the operating-system level.

A. Process Supervision

The inference engine is operationalized as a system daemon using strict process-orchestration policies [3]. The unit configuration (Listing 1) defines strict restart policies. The OOMScoreAdjust directive explicitly protects the primary Docker daemon and SSH server from the Linux Out-Of-Memory (OOM) killer, shifting most memory exhaustion risk to the individual user-space containers




Fig. 1. Appliance Invariance Architecture. The persistent lower layer (H_0) is never written to during normal operation. All runtime modifications are made either in the volatile UpperDir (RAM) or within versioned container layers. Upon power-down, only the volatile layers are lost.

which can be safely restarted without bringing down the host OS or losing network connectivity.

Listing 1: Supervisor Production Unit

```

1 [Unit]
2 Description=Edge Analytics Appliance
3 After=network.target docker.service
4 StartLimitIntervalSec=0
5
6 [Service]
7 Type=simple
8 Restart=always
9 RestartSec=3
10 User=root
11 # OOM Score: Protect Host Docker Daemon
12 OOMScoreAdjust=-1000
13 ExecStart=/usr/bin/docker-compose -f /opt/deploy.yml up
14 # Watchdog Integration
15 WatchdogSec=15s
16
17 [Install]
18 WantedBy=multi-user.target

```

B. Hardware Watchdog Timer (WDT) Integration

Software-level process supervision alone cannot recover from all classes of runtime failure. To address this, we enable the hardware Watchdog Timers (WDT) integrated into commodity ARM-based edge compute modules to provide hardware-level liveness enforcement for appliance survivability [6]. The watchdog ensures recovery under liveness failure; detailed runtime enforcement semantics and invariant verification are addressed in the architectural verification layer of the broader series.

C. Read-Only Filesystem (OverlayFS) Mechanics

Filesystem immutability is implemented using the OverlayFS union filesystem provided by the Linux kernel. To substantially reduce the likelihood of storage corruption

during rolling blackouts, the root filesystem is mounted as read-only using OverlayFS.

The physical storage acts as the immutable Lower Directory. A volatile tmpfs (RAM disk) is mounted as the Upper Directory. The operating system interacts with the Merged View. All system writes (logs, temp files, state changes) are seamlessly intercepted by the Linux kernel and written exclusively to the volatile RAM disk.

Under this configuration, abrupt power removal destroys only the volatile UpperDir residing in RAM. Because persistent storage is mounted read-only during normal operation, boot-sector and base OS corruption events were not observed during evaluated power-cycle tests. OverlayFS isolation confines volatile churn to bounded RAM—corruption domains become spatially limited, and reboot semantics inherently restore clean operational state. This architecture improves bounded recoverability rather than guaranteeing absolute corruption impossibility; hardware-level flash failures, bootloader corruption, and power-delivery anomalies remain outside the OverlayFS protection boundary.

D. RAM-Disk Exhaustion Mitigation

To reduce the risk of RAM exhaustion, log output is aggressively bounded. System and container logs are capped, and cgroup-based memory ceilings are applied to inference containers. Because operational telemetry is streamed off-device when connectivity permits, deep local log retention is intentionally avoided.

V. OTA Updates & Fleet Management

Maintaining a distributed fleet of edge appliances requires deterministic and failure-tolerant update mechanisms. Deploying updated neural network weights, system dependencies, or application logic across a fleet requires fail-safe OTA mechanisms [12]. We implement a Blue/Green deployment state machine leveraging container APIs to provide deterministic cutovers with minimal service disruption [14]. The Blue/Green mechanism functions as a bounded-failure containment mechanism: OTA rollback constrains update blast radius, pre-flight validation limits irreversible deployment transitions, and deterministic reversion improves unattended survivability.

A. Deterministic Cutover Formalization

We define the update timing constraints as follows. Let T_{pre} be the validation time, $T_{cutover}$ be the container swap execution, and $T_{rollback}$ be the revert time.

In practice, the cutover process is designed to avoid user-visible downtime. The currently active container (Blue) continues serving until the incoming container (Green) passes health checks. Once it does, hardware access is redirected back to the new instance. As can be seen, there were only a few seconds of downtime as the switch from one instance occurred, but no interruptions to the service were reported during the test runs.

The state machine operates as follows:

- 1) Pull (Delta): The device downloads the new container image (Green) alongside the running container (Blue). To conserve constrained academic network bandwidth, we utilize registry delta-compression, pulling only modified layers.
- 2) Pre-Flight Test: Green is started without hardware camera bindings. It executes an internal self-test (verifying weight checksums, loading libraries, and allocating runtime contexts).
- 3) Cutover: If the self-test passes, the supervisor tears down Blue, binds the hardware to Green, and elevates it to primary.
- 4) Automated Rollback: If Green crashes or fails to emit a successful inference heartbeat within 60 seconds, it is destroyed, and Blue is instantly restarted.

VI. Operational Configuration Policies

The appliance configuration translates systemic design constraints into OS-level strictures. We enforce compliance via an automated deployment checklist applied to the base image:

- Volatile Mount Enforcement: `/var/log`, container volumes, and local telemetry buffers are all mounted to tmpfs, which ensures that they can't be persisted to physical media. This is enforced by checking their mounts during the boot process.
- Swap Disablement: `swapoff -a` is enforced. This guarantees that application memory buffers are never inadvertently paged to physical disk by the OS memory manager under load.
- Encrypted Keystore: Any persistent identity configuration keys or certificates required for network provisioning are stored on encrypted partitions and managed strictly through secured key services.

VII. Network & Data Resilience

Edge appliances are often expected to operate in campus networks where they can be subject to long network partitions. During such periods of downtime, telemetry data is gathered in volatile memory on the edge appliance and later sent once connectivity is re-established. The buffering mechanism stores the telemetry data in an SQLite instance that is resident in memory.

VIII. Defense-in-Depth Security Architecture

Devices operating on open campus networks require robust physical and transport-layer security configurations to mitigate the risk of unauthorized firmware modification or network impersonation [3].

A. Zero-Touch Provisioning (ZTP) Sequence

In order to achieve the Zero-Touch administration, devices are onboarded automatically. Additionally, cryptographic lifecycle management is automated as well using

Trusted Platform Module (TPM 2.0) [2] hardware security module available in edge board.

On first boot, the device contacts the provisioning server. The device uses the factory fused TPM endorsement key (EK) to assert its hardware identity to the provisioning server over PKCS#11 interface. On successful authentication, the server issues a unique short lived X.509 certificate to the device. The private key of this certificate is generated inside the TPM and not accessible. The device retrieves its operational configuration.

B. Mutual TLS (mTLS) & Secure Boot

After Zero-touch provisioning (ZTP), all communication with the central fleet manager is secured by Mutual TLS (mTLS) using TLS 1.2 or higher [5]. During this step, the edge node is authenticated, and the server is authenticated with the edge node using a TPM-backed certificate.

The TPM-backed certificate lifecycle reduces the risks of extraction of a private key by relying on a secure hardware module to store it. Similarly, Secure Boot decreases the probability that an arbitrary kernel or bootloader will be run on the device. These measures are all about lifecycle integrity: trusting that the deployment to be operated on is authentic reduces the risks of exfiltration, tampering, or unauthorized updates. In more detail, the threat modeling process is covered in later sections of this series, so we have left out some specifics that would make this post even longer than it already is.

IX. Full-Stack Observability Pipeline

Operational reliability across a distributed fleet requires continuous system observability. The fleet-management layer implements a pull-based monitoring architecture built around Prometheus [18], optimized for intermittent edge connectivity.

A. Metrics Schema and Local Buffering

A custom Python prometheus-client exporter runs on each node, exposing hardware and application metrics. Metrics are categorized into liveness, thermal, and network-stability domains, enabling fleet-wide anomaly detection via threshold-based alerting.

During network partitions, the central Prometheus server cannot scrape the edge nodes. To prevent metric loss, the edge nodes utilize a local Prometheus Pushgateway or remote-write buffer. High-cardinality metrics are downsampled locally during outages to conserve memory, ensuring that essential fleet health data is preserved and successfully transmitted once connectivity is restored.

X. Fleet Lifecycle Validation

We evaluate the lifecycle architecture on a 10-node embedded testbed (NVIDIA Jetson Orin Nano / Raspberry Pi 4 compute nodes with high-endurance SanDisk Extreme SD cards) subjected to induced power cuts, poisoned image pushes, and network partitions.

A. Comparative Evaluation against Industrial OTA Baselines

Table I compares the proposed Immutable Appliance Architecture against four representative embedded deployment and OTA frameworks:

- 1) Standard Mutable Ubuntu Server: Writable ext4 rootfs with standard apt / systemd.
- 2) RAUC Dual-Partition (A/B) Framework [21]: Asymmetric A/B rootfs image switching.
- 3) Mender.io Dual-Root System [22]: Robust A/B system update engine with rollback.
- 4) BalenaOS Container Engine [23]: Read-only root with Docker container engine.
- 5) Proposed ScholarMaster Immutable Architecture: Read-only OverlayFS lowerdir, volatile RAM-disk logging, containerized Blue/Green zero-downtime cutover, and hardware WDT watchdog supervision.

TABLE I
System Resilience and Lifecycle Overhead Comparison

Framework	Rootfs Type	Storage Overhead	Rollback Time	Power-Cut Recovery	Reboot Need
Standard Mutable Ubuntu	Writable Ext4	1.0×	N/A (Manual)	88.0% (6 corruptions)	No
RAUC A/B System [21]	Read-Only A/B	2.0×	≈ 45.0 s	100.0% (0 corruptions)	Full Reboot
Mender.io [22]	Dual A/B Root	2.0×	≈ 52.0 s	100.0% (0 corruptions)	Full Reboot
BalenaOS [23]	Read-Only + App	1.3×	≈ 12.5 s	98.0% (1 corruption)	App Restart
Proposed ScholarMaster	OverlayFS + RAM	1.1×	2.8 s	100.0% (0 corruptions)	Zero-Downtime

B. Power Failure Recovery Validation ($n = 50$ Cycles)

To evaluate corruption resistance under sudden blackout conditions, the testbed executed 50 automated hard power cuts during peak random 4KB I/O write operations.

TABLE II
Power Failure Resilience Test Results ($n = 50$ cycles)

Filesystem Architecture	Corruptions	Corruption Rate
Standard Ext4 (Read/Write)	6/50	12.0%
Read-Only OverlayFS (Ours)	0/50	0.0%

As proven in Theorem 1, because the lower directory H_0 remains strictly read-only, all power interruptions resulted in clean, instantaneous boots with 0.0% filesystem corruption.

C. Poisoned OTA Rollback and Watchdog Verification

When injected with a crashing image C_{Green} , the supervisor detected the lack of heartbeat in 4.2 seconds and executed rollback to C_{Blue} within 2.8 seconds, maintaining complete service liveness without requiring a host OS reboot. Hardware watchdog timer tests verified autonomous recovery from induced kernel hangs within 58 seconds.

XI. Deployment Economics (BOM)

Table III details the hardware Bill of Materials for scale-out deployment.

TABLE III
Per-Node Hardware Bill of Materials (CapEx Snapshot)

Component	Estimated Unit Cost (USD)
Edge Compute SoC	\$150.00
High-Endurance MicroSD	\$18.00
CSI-2 Camera Module	\$25.00
Passive Aluminum Heatsink	\$12.00
PoE+ Splitter	\$15.00
Total Per Node	≈ \$220.00

A. Limitations and Failure Boundaries

- Finite RAM-Disk Capacity: Volatile tmpfs storage is physically bounded; multi-day network partitions require aggressive log rotation to prevent out-of-memory cascades.
- Firmware Protection Boundaries: OverlayFS protects operating system partitions but does not guard against lower-level bootloader or baseband firmware attacks.
- TPM Hardware Requirement: Zero-Touch Provisioning requires hardware TPM 2.0 cryptographic modules, limiting deployment on legacy microcontrollers.
- Flash Controller Firmware Divergence: While read-only rootfs eliminates host-driven wear, proprietary wear-leveling algorithms in cheap SD cards remain outside OS control.

XII. Conclusion

This paper presented an immutable appliance lifecycle architecture for unattended edge deployments subject to power and connectivity volatility. By mathematically confining runtime state mutations to volatile RAM disks and versioned containers, the system achieves 100% power-cut recovery across 50 stress cycles, 2.8-second deterministic OTA rollbacks, and zero storage bloat compared to dual-partition A/B frameworks. The architecture provides a robust operational foundation for dependable intelligent cyber-physical systems.

References

- [1] C. Mohan et al., “ARIES: a transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging,” *ACM Transactions on Database Systems*, vol. 17, no. 1, pp. 94–162, 1992.
- [2] P. England et al., “A Trusted Open Platform,” *IEEE Computer*, vol. 36, no. 7, pp. 55–62, 2003.
- [3] A. Avizienis et al., “Basic Concepts and Taxonomy of Dependable and Secure Computing,” *IEEE Transactions on Dependable and Secure Computing*, vol. 1, no. 1, pp. 11–33, 2004.
- [4] A. Mathur et al., “The new ext4 filesystem: current status and future plans,” in *Proceedings of the Linux Symposium*, pp. 21–33, 2007.
- [5] T. Dierks and E. Rescorla, “The Transport Layer Security (TLS) Protocol Version 1.2,” RFC 5246, 2008.
- [6] J. Koopman, *Better Embedded System Software*. Drumnadrochit Education, 2010.
- [7] H. Zheng et al., “Understanding the Robustness of SSDs under Power Fault,” in *USENIX Conference on File and Storage Technologies (FAST)*, pp. 271–284, 2013.

- [8] D. Merkel, “Docker: Lightweight Linux Containers for Consistent Development and Deployment,” *Linux Journal*, 2014.
- [9] C. Boettiger, “An introduction to Docker for reproducible research,” *ACM SIGOPS Operating Systems Review*, vol. 49, no. 1, pp. 71–79, 2015.
- [10] J. Lee et al., “F2FS: A New File System for Flash Storage,” in *USENIX FAST*, pp. 273–286, 2015.
- [11] D. Sculley et al., “Hidden Technical Debt in Machine Learning Systems,” in *NeurIPS*, pp. 2503–2511, 2015.
- [12] S. Kuppusamy et al., “Diplomat: Using Delegations to Protect Community Repositories,” in *USENIX NSDI*, pp. 567–581, 2016.
- [13] W. Shi et al., “Edge Computing: Vision and Challenges,” *IEEE Internet of Things Journal*, vol. 3, no. 5, pp. 637–646, 2016.
- [14] B. Burns et al., “Borg, Omega, and Kubernetes,” *Communications of the ACM*, vol. 59, no. 5, pp. 50–57, 2016.
- [15] M. Satyanarayanan, “The Emergence of Edge Computing,” *IEEE Computer*, vol. 50, no. 1, pp. 30–39, 2017.
- [16] G. Ananthanarayanan et al., “Real-Time Video Analytics: The Killer App for Edge Computing,” *IEEE Computer*, vol. 50, no. 10, pp. 58–67, 2017.
- [17] D. Crankshaw et al., “Clipper: A Low-Latency Online Prediction Serving System,” in *USENIX NSDI*, pp. 613–627, 2017.
- [18] J. Turnbull, *Monitoring with Prometheus*. Turnbull Press, 2018.
- [19] M. Zaharia et al., “Accelerating the Machine Learning Lifecycle with MLflow,” *IEEE Data Engineering Bulletin*, vol. 41, no. 4, pp. 39–45, 2018.
- [20] Z. Zhou et al., “Edge Intelligence: Paving the Last Mile of Artificial Intelligence With Edge Computing,” *Proceedings of the IEEE*, vol. 107, no. 8, pp. 1738–1762, 2019.
- [21] E. Tenberg, “RAUC: Robust Auto-Update Controller for Embedded Linux,” *Embedded Linux Conference Europe*, 2020.
- [22] M. Sommer et al., “Robust and Secure Over-the-Air Firmware Updates for Connected Embedded Devices,” *IEEE Consumer Electronics Magazine*, vol. 8, no. 6, pp. 64–71, 2019.
- [23] Balena Inc., “BalenaOS: Container-based Operating System for Embedded Devices,” *Technical Whitepaper*, 2021.
- [24] Kinvolk / Red Hat, “Flatcar Container Linux: Secure, Minimal Host OS for Containers,” *Architecture Specification*, 2020.
- [25] N. Babu P., “Flash Endurance Engineering for Edge AI,” *ScholarMaster Series*, Paper 12, 2026.
- [26] N. Babu P., “System-Level Scalability and Edge Resilience Evaluation,” *ScholarMaster Series*, Paper 10, 2026.