

Beyond the Panopticon: A Longitudinal Study of Student Trust in Automated Stewardship Systems

Dr. S. Suresh Kumar
Principal

Swarnandhra College of Engineering & Technology (Autonomous)
Seetharampuram, Narsapur, Andhra Pradesh, India

Abstract—Automated monitoring systems, which are pervasive in modern educational institutions, may cause situations in which student privacy is at odds with efficiency. Students may simultaneously desire the convenience of automated services, and the lack of automated sensing.

This study examines whether the visible enforcement of privacy constraints influences how trust develops within institutional monitoring environments. A longitudinal field study ($N = 540$) was conducted across a sixteen-week academic semester consisting of two deployment phases. During the initial phase the system operated in a conventional opaque monitoring configuration. In the second phase a series of transparency interventions was introduced, including skeletal abstraction displays and student-accessible audit dashboards.

Following the transparency interventions, mean Trust Scores increased from 2.4/5.0 to 4.1/5.0 on a validated System Trust Scale, while reported surveillance anxiety declined substantially. A strong positive correlation ($r = 0.82$) was observed between exposure to visible abstraction mechanisms and acceptance of the monitoring environment.

The results indicate that privacy safeguards become socially meaningful only when their operational boundaries are perceptible to the individuals affected by the system.

Index Terms—Automated Stewardship, Visible Abstraction, Transparency-Mediated Trust, Sociotechnical Acceptance, Contextual Integrity, Longitudinal Sentiment Analysis.

I. Introduction

The dominant thesis of this paper is that visible operational constraints materially influence institutional trust formation in monitoring environments. Monitoring technologies increasingly operate as infrastructural components within modern educational environments, automating attendance tracking, classroom analytics, and safety monitoring through sensor networks that remain largely invisible to the individuals they observe.

A. Why Hidden Safeguards Fail to Generate Trust

Modern monitoring architectures often incorporate extensive privacy-preserving mechanisms, including data minimization pipelines and automated deletion policies. However, backend privacy protections that remain invisible to users have limited social relevance. Users cannot psychologically evaluate safeguards they cannot perceive. Operational opacity produces worst-case assumptions: when confronted with opaque infrastructure, students often assume that their physical movements, facial expressions, and private interactions are continuously recorded,

analyzed, and permanently archived. This perception generates what surveillance scholars describe as a Chilling Effect [14]. Visible boundedness changes perceived institutional intent—trust depends on perceptibility, not merely technical correctness.

B. Study Framing and Subsystem Scope

This study adopts the framing of Automated Stewardship [6], which positions monitoring technologies as bounded systems designed to serve specific, transparent institutional purposes while visibly demonstrating their own limitations. We designed a semester-long longitudinal field study to evaluate whether visible transparency mechanisms—such as real-time skeletal abstraction displays and user-accessible audit dashboards—can alter student perceptions of monitoring technologies. The intervention focused purely on making existing operational constraints perceptible to the student body.

Within the ScholarMaster architecture, this subsystem studies visible privacy, transparency-mediated trust, and sociotechnical acceptance only. It does not define irreversible architecture (Paper 17), validate runtime enforcement (Paper 18), derive formal verification theorems (Paper 21), synthesize deployment infrastructure (Paper 20), or optimize federated coordination (Paper 14).

C. Data Availability

The anonymized survey instruments (questionnaires), interview protocols, and aggregated Likert datasets used in this sociological study are available from the authors upon reasonable request to support reproducibility in HCI research.

II. Theoretical Framework

Understanding this relationship requires examining established theories within human–computer interaction and privacy research.

A. The Privacy Paradox

Within Human–Computer Interaction research, the Privacy Paradox is the observed discrepancy between users’ publicly stated concerns about privacy and their actual behaviors when interacting with digital technologies. In other words, people tend to state that their privacy is

critically important to them while regularly giving away personal information for only minor benefits [17].

Kokolakis’ review of the existing literature suggests that this paradox may stem from an insufficient understanding of what happens with the data after collection. In other words, users who do not have a clear idea of what happens to their data can dramatically underestimate its relevance at the time of providing it and express concerns in the aftermath of a particular data processing step being completed.

In educational environments there is a related but different dynamic. Students are often willing to adopt technologies that are framed as helpful (or even mandatory) so long as they are not explicitly presented as forms of surveillance. One possible explanation for this is that it is not the sensor technology itself that determines acceptance by users, but rather the agency and transparency of the system operators and processes around the technology.

B. Contextual Integrity

The expectations of privacy are dependent on the contexts regarding information-flow. Nissenbaum’s notion of contextual integrity (CI) provides a framework for representing and analyzing informational privacy values in complex domains [7]. In particular, privacy violations are characterized by information flows that deviate from informational norms in a particular context.

For example, keeping track of who has attended the class in a classroom setting is par for the course - students know they are being monitored. By contrast, performing an emotional analysis of them for commercial reasons would most definitely be surprising. A monitoring technology that will tell the user what it can and cannot do seems to be more intuitive than a monitoring device that will do whatever it wants.

C. Surveillance Capitalism in Education

Today’s online systems have been increasingly leveraging data about users’ behavioral patterns as an input for analytical processes [21]. Learning analytics platforms represent one such system that has been using data about students’ interactions with the system to evaluate their level of engagement and emotional responses. Slade and Prinsloo [8] express concern regarding the tendency of learning analytics to depersonalize student behavior, thus undermining their privacy and ignoring the ethical implications of such surveillance. The authors suggest that privacy-oriented learning analytics should be designed to postulate the informational accountability of its stakeholders by being transparent about its inner workings, thus adhering to the Privacy by Design philosophy proposed by Cavoukian [6]. By doing that, educational institutions would promote student engagement by fostering a trusting environment in which both parties are accountable for their actions.

D. Trust in Automation

Trust in automated systems is typically modeled as a function of predictability, dependability, and perceived system intent [4].

Opaque systems naturally undermine predictability because users simply cannot observe how the system behaves internally or what triggers its decisions. Mayer’s highly influential trust model identifies three key antecedents to organizational trust: Ability, Benevolence, and Integrity [3].

Within automated monitoring contexts, we map these antecedents directly to system features:

- Integrity is supported through user-verifiable auditability (e.g., viewing one’s own data logs).
- Benevolence can be communicated through non-invasive abstraction mechanisms (e.g., ensuring cameras cannot see faces).

Transparency mechanisms therefore operate not only as technical safety guardrails but also as vital psychological signals that shape the formation of trust.

III. Methodology

To evaluate the influence of visible privacy mechanisms, we conducted a semester-long longitudinal observational study within an institutional monitoring deployment. The sensing and transparency interfaces were deployed within a controlled institutional testbed environment configured to emulate a campus-wide monitoring system; importantly, the intervention targeted interface transparency rather than changes to sensing hardware. Although the study was conducted within a single institution, the mechanisms of visible transparency and contextual integrity align closely with established trust literature.

A. Participants

The study cohort consisted of $N = 540$ undergraduate students drawn from computer science and electronics engineering disciplines.

TABLE I
Participant Demographics Detail

Category	Count (N=540)	Percentage
Gender		
Male	310	57.4%
Female	230	42.6%
Year of Study		
2nd Year (Sophomore)	150	27.8%
3rd Year (Junior)	200	37.0%
4th Year (Senior)	190	35.2%

- Tech Literacy: High (Engineering majors).
- Participation: Participation remained strictly voluntary and was explicitly separated from academic grading or evaluation to prevent coercion.

B. Experimental Phases

The academic semester was divided into two sequential phases to isolate the effect of transparency interventions (one baseline phase and one transparency-augmented phase).

Phase 1 (Weeks 1 – 8): The surveillance system was running on a standard, opaque mode. The students knew that they were under observation through the typical signs in the corridors, but the details of the system remained unknown to them. Overall, it did not look any different from regular CCTV cameras.

Phase 2 (Weeks 9–16): A coordinated series of transparency interventions was introduced into the environment:

- **Skeleton Abstraction Displays:** Hallway monitors were activated to show the system’s actual view: abstract stick-figures rather than RGB video.
- **Audit Dashboards:** A secure mobile application allowed students to view their own anonymized event logs.
- **Data Deletion Mechanisms:** Students were granted the ability to purge their own temporal records.
- **Ambient LED Indicators:** Cameras were fitted with LED rings indicating real-time active/inactive processing states.

Crucially, the underlying sensing hardware and analytics pipeline remained completely unchanged between phases. Only the visibility of the system’s operations differed.

C. Survey Instrument Validation

The primary measurement tool, the System Trust Scale (STS), was administered at Week 8 (the conclusion of the opaque phase) and Week 16 (the conclusion of the transparent phase).

Before using the instrument, it was pilot tested with a small group of participants ($n = 50$). The reliability of the item was found to be remarkably high. Specifically, Cronbach’s Alpha was $\alpha = 0.88$ for Trust and $\alpha = 0.91$ for Anxiety. Both figures are well above the acceptable threshold of 0.70 [1].

TABLE II
System Trust Scale (STS) Instrument Items

Construct	Item Phrasing (5-Point Likert)
Surveillance	“I feel watched/judged by the system.” “I worry about who can see my data.”
Utility	“The system improves my campus experience.” “Automated attendance saves valuable class time.”
Trust	“I trust that my personal data is safe.” “I believe the system works as described.”
Agency	“I feel in control of my data.” “I understand how to delete my records.”

D. Internal Validity Considerations

While the study provides an excellent basis for comparison, based on its longitudinal nature, there are a few factors that may affect the result’s internal validity. First, it is possible that over 16 weeks, the participants grew accustomed to the hardware, thus reducing their anxiety naturally. Second, positive feedback could have been prompted by the novelty effect, as Phase 2 used modern equipment with improved UI/UX. Finally, social desirability bias always affects survey-based research, and this study is not an exception.

IV. Results

The transition between opaque and transparent deployment phases produced substantial shifts in reported student sentiment.

A. Quantitative Sentiment Shift

The transition from Phase 1 to Phase 2 produced substantial shifts in student sentiment across all measured categories, as detailed in Table III.

TABLE III
Mean Student Sentiment Scores (1=Negative, 5=Positive)

Metric	Phase 1	Phase 2	Delta (Δ)
Trust in System	2.4 ± 1.1	4.1 ± 0.8	+1.7
Anxiety Level	4.2 ± 0.9	1.8 ± 0.7	-2.4
Perceived Utility	3.1 ± 1.2	4.5 ± 0.6	+1.4
Understanding of Data	1.5 ± 0.5	4.8 ± 0.4	+3.3

Mean scores for Trust increased considerably from 2.4 to 4.1, whereas Anxiety dropped significantly from 4.2 to 1.8. Understanding of Data improved the most, jumping by 3.3 points. As can be seen from phase 1, students were unsure what information exactly was being gathered about them. It can be argued that seeing the skeleton-like image of the face helped alleviate their concerns that their personal details are not being stored, hence why the score of Anxiety diminished.

B. Demographic Variance

Sentiment changes were not uniform across participant groups.

Disciplinary Differences: On the initial measurement, Computer Science students reported significantly lower levels of trust compared to other engineering disciplines (2.1 versus 2.6). This could be explained by their superior understanding of the technological possibilities of continuous video surveillance. On the other hand, they also demonstrated a larger increase in trust after receiving the transparency treatment (Delta +1.9 versus +1.5). This suggests that tech-savvy users are especially receptive to truth-telling transparency treatments. After all, they quickly realized that the allegedly restricted system was indeed physically limited and thus could restore their trust by becoming aware of these restrictions.

Gender Dynamics: Female participants scored higher on Anxiety in Phase 1 (4.6 versus 3.9) possibly due to general social trends towards increased surveillance. This discrepancy was less prominent in Phase 2 (1.9 versus 1.7), possibly because users’ perception of privacy is affected more by visible signs of monitoring.

C. Statistical Significance Testing

Prior to parametric testing, assumption checks for normality (Shapiro-Wilk) were conducted. We performed a paired sample t-test comparing Phase 1 and Phase 2 scores for each construct. The observed improvements in Trust ($t(539) = 14.2, p < 0.001$) and Understanding ($t(539) = 21.5, p < 0.001$) met conventional statistical significance thresholds.

Furthermore, a Cohen’s d effect size analysis yielded $d = 1.8$ (95% CI [1.62, 1.98]) for the Understanding of Data metric, consistent with a large effect size associated with the transparency intervention. A regression-based mediation analysis (Indirect effect = 0.41, 95% CI [0.32, 0.52]) suggests that reduced surveillance anxiety partially mediates the relationship between visual transparency and trust formation. The mediation analysis suggests influence rather than universal causality; cohort composition affects interpretability. Notably, 7% of participants showed no change or slight decreases in Trust scores, suggesting heterogeneous adaptation patterns within the evaluated cohort.

D. The Role of Visual Privacy Cues

Survey responses revealed an interesting hierarchy regarding which specific features actually drove the shift in perception. When asked to identify the most influential factor, 82% of participants pointed to the Skeleton View visualization. In sharp contrast, the complex backend technical mechanisms were cited by only 15% of students.

This disparity highlights an important design implication: perceptible privacy signals—interface-mediated constraint visibility and user-observable abstraction—carry vastly greater psychological weight than invisible technical guarantees. Visibility is not merely aesthetic; it functions as a trust-regulation mechanism. Users must be able to visually infer the system’s operational limits to feel comfortable with its presence.

E. Qualitative Feedback Clusters

We conducted exit interviews with 50 randomly selected participants to gather qualitative insights. Thematic mapping was conducted using inductive thematic coding with an inter-rater reliability of Kappa = 0.85. Thematic analysis (Table IV) reveals a conversational shift from fear of judgment to an appreciation of utility.

This qualitative shift suggests that the perceived nature of the relationship changed. The system was viewed less as an external authority and more as a tool under the student’s partial purview. While most participants




Fig. 1. The Skeleton View (Visible Anonymity) was overwhelmingly the most frequently cited driver of trust. Backend technical features were reported as significantly less impactful on sentiment than visible frontend cues.

TABLE IV
Qualitative Theme Mapping: Before vs. After

Phase 1 Theme	Phase 2 Theme
“It feels like big brother is watching.”	“It’s just checking if I’m raising my hand.”
“What if it records me sleeping?”	“I saw the skeleton view; it doesn’t see my face.”
“The data will be used against me.”	“I can delete my log anytime, so I don’t worry.”

described increased reassurance, 3 interviewees expressed continued discomfort with any form of automated monitoring regardless of abstraction.

F. Data Deletion Usage

During Phase 2, a mechanism allowing students to delete their generated metadata was made available.

- Unique Users: 43 (8% of population).
- Follow-up Behavior: 38 of the 43 students subsequently re-engaged with the system.

This behavior suggests that perceived agency may contribute more to trust formation than the actual frequency of data removal. Knowing they possessed the capability to delete their data appeared to provide students with the psychological safety required to participate.

G. Correlation Analysis

We calculated the Pearson correlation coefficient between visualizability of privacy (measured via binary exposure to the intervention) and acceptance of monitoring (measured via the STS Trust sub-scale). The resulting $r = 0.82$ indicates a strong positive correlation, suggesting a robust association between the deployment of visible privacy mechanisms and the successful formation of institutional trust.

V. Trust Mechanics Analysis

To interpret the observed sentiment shifts, we formalize a mediated trust-formation model grounded in established

organizational trust theory. We introduce an explicit mediation model grounded in Mayer et al.’s foundational definitions of Cognitive Trust (rational assessment of ability and integrity) and Affective Trust (emotional security and benevolence) [3].

A. A Formal Trust Mediation Model

Let the primary variables be defined as:

- V = Visual Privacy Cues
- A = Affective Trust
- C = Cognitive Trust
- S = System Acceptance

The following functional relationships mapping transparency to final acceptance are proposed:

$$A = f_1(V) \quad (1)$$

$$C = f_2(A, V) \quad (2)$$

$$S = f_3(C) \quad (3)$$

The observed data are consistent with a mediated trust formation hypothesis in which visible abstraction reduces affective anxiety ($A = f_1(V)$), which subsequently enables cognitive verification ($C = f_2(A, V)$), ultimately leading to system acceptance ($S = f_3(C)$). Critically, the temporal progression matters: affective trust appears to function as a necessary precursor for cognitive trust. This model should be understood as a sociotechnical mediation framework rather than a universal causal law; the associations are strong within the evaluated cohort but may vary across populations and institutional cultures.

B. Cognitive vs. Affective Dynamics

Building upon Mayer’s proposition [3], the student dashboards were primarily focused on Cognitive Trust. In this case, students could easily justify that the information contained in the metadata was limited. At the same time, having access to this data made them realize that they were not withholding any private information from the system operators.

On the other hand, the visual abstraction overlays and ambient indicators directly addressed Affective Trust. This consideration was aimed at the user’s fear of being watched. The visual indication of the camera’s gaze as a skeletal abstraction helped to allay these fears by “showing” the user that the machine was not out to get them.

Critically, it is essential to note that the temporal progression is vital, as affective trust seems to act as a precursor to cognitive trust. The model should be regarded as a sociotechnical mediation framework, not a universal law, as the associations were found to be strong within the studied sample, but they could differ in other populations and institutional settings.




Fig. 2. The Trust Loop Model. Inspired by Mayer et al., [3], the figure abstracts visual content based on which a user can develop emotional safety (Affective Trust) that leads to rational auditing behaviors (Cognitive Trust), which finally determines user acceptance.

VI. Discussion

The study highlights how transparency mechanisms reshape the relationship between institutional monitoring systems and their participants.

A. From Asymmetry to Symmetry

Traditional surveillance systems are characterized by a specific paradigm in which the observer’s identity is hidden while the observed is fully exposed [2]. Such one-sided information availability may lead to such phenomena as mistrust in an institution. The proposed intervention strategy uses the reversed version of surveillance in which all collected data are available to the observed individual while his or her personal identity is hidden for the institution at the surveillance system side. Such an approach provides a better balance of power in the relationship between the student and the institution.

B. Algorithmic Fairness and Bias Perception

A significant element of trust is the feeling of fairness. Both Benjamin [22] and Noble [18] have demonstrated that automated systems often contain and perpetuate hidden demographic biases through seemingly neutral optimization. In our study, minority students were initially more suspicious in phase 1, and their concerns were related to the possible unfairness of the presented algorithm in facial recognition technology.

However, the skeletal abstraction display seemed to directly address these issues. By visually depicting the system’s tracking of geometric positions rather than attempting to identify physical characteristics, skin color, or cultural dress, it demonstrated its impartiality. Essentially, the display served as a continually updated




Fig. 3. Information Asymmetry and Symmetry. In (a) only the administrator can see, while in (b) both sides share the information from the Transparency Dashboard, thus establishing trust.

”model card” [23], illustrating the system’s functioning and parameters.

C. Jurisdictional Context

Although the GDPR [13] is an excellent benchmark for privacy research, it is essential to recognize that regulations may differ at the national level. For instance, the recently proposed Digital Personal Data Protection (DPDP) Act in India focuses on the obligations of the Data Fiduciaries. Systems that reduce the storage of biometric data and emphasize abstracting away would alleviate the due diligence of such organizations under the proposed legislation. By ensuring that no raw biometric data is stored, the burden of liability is significantly reduced, which would encourage adoption by institutions.

D. The Role of Explainability

A more complex algorithm is often suspected by the end-users. This became evident when the inference logic was made public on the dashboard, and the number of complaints dropped significantly. It is essential to remember that transparency of data flow and inference logic are not the same, and users need both. It was evident in the case, as users needed the clarity of the flow of data (which was provided by the mobile dashboards). On the other hand, the visualized skeletal logic showed how the observations were interpreted and contributed to the final results. Both aspects of explainable AI increased the level of trust between users and the system.

E. The Paradox of Transparency

Transparency itself is a design problem. The initial versions of the student dashboard exposed the raw JSON logs from the server, which proved to be too technical

for most students to understand. Worse, this complexity seeded mistrust since users felt they had no control over the information. Thus, our early designs failed the ”transparency test” and needed to be redone. These experiences taught us an important lesson about transparency design: it should remain intelligible to avoid mistrust and cognitive overload.

F. Mitigating the Chilling Effect

Surveillance literature discusses the concept of Chilling Effect, which is manifested in the diminishment of the proper behavior due to the sensed or perceived risk of being observed [14]. In Phase 1, we could see signs of such fear among students who tried to make their bodies less ”suspicious” in front of the cameras. In Phase 2, we saw that such fears were to a certain extent diminished, as there was less ”positional tension” in the students’ posture. The visual abstraction techniques helped to ensure that the students’ posture or other behavioral pattern would not be used against them.

G. Institutional Implications

Administrators’ initial worries about transparency leading to gaming the system appear to be misplaced, based on the results. The study even shows the opposite trend: that transparency promotes cooperative attitudes. Students who learned that facial video recording was not happening were less likely to try to hide from or evade the cameras.

H. Limitations

The sociotechnical assessment of the visible privacy bounds offers several limitations to the study:

- Novelty effects: The initial enthusiasm for the clean mobile dashboard introduced in the second phase may have contributed to an upward bias in satisfaction and positive attitudes toward the monitoring system.
- Self-selection: Students with preexisting privacy concerns were more likely to participate in the study, skewing the results toward more privacy-conscious individuals.
- Hawthorne effects: The awareness of being observed and studied may have influenced students’ attitudes toward the monitoring software.
- Single-campus cultural bias: Institutional culture can shape trust-building strategies; thus, the visible privacy bounds may have different effects in different organizational contexts.
- Age and occupation skew: The sample was biased toward engineering students who were more tech-savvy and had a better understanding of the abstracted concepts.
- Self-reported Likert scales: Longitudinal sentiment analysis relied on self-reported attitudes toward the system, which were subject to biases and social desirability.

- Longitudinal Uncertainty: While a 16-week study may well capture initial adoption reaction patterns, it cannot assess long term acceptance by subsequent generations within an extended multi-generational cohort. Institutionalization, shifts in perception, and changing novelty over time may well outweigh positive first impressions within the confines of the semester.

VII. Conclusion

This longitudinal intervention study explored the role that visible transparency mechanisms could play in fostering trust in the context of an educational data monitoring system. After transitioning from an initially opaque implementation to a skeletal abstraction display plus student-accessible audit dashboards, mean Trust scores in the intervention cohort rose from 2.4 to 4.1 (out of 5) during the study period.

The results imply that privacy-preserving backend mechanics, such as encryption and deletion policies, have limited power to shape user perception if they are not directly presented to the consumer. Instead, the study suggests that presenting constraints to the consumers through interface-mediated legibility encourages the emergence of trusting attitudes stemming from the users' awareness of the operator's agency and accountability. The trust generation process appears to be mediated through a mechanism which involves reassurance conveyed by verification means followed by behavioral acceptance but is limited by situational factors and requires further investigation.

Within the ScholarMaster architecture, this framework only provides the visible privacy and institutional trust subsystem, studying the perceptible system boundedness, interface-mediated constraint visibility, and sociotechnical acceptance, without considering the cryptographic ledgering, edge intelligence, architectural irreversibility doctrine, or spatiotemporal verification rules.

References

- [1] L. J. Cronbach, "Coefficient alpha and the internal structure of tests," *Psychometrika*, vol. 16, no. 3, pp. 297–334, 1951.
- [2] J. Bentham, *The Panopticon Writings*. Verso, 1995.
- [3] R. C. Mayer, J. H. Davis, and F. D. Schoorman, "An integrative model of organizational trust," *Academy of Management Review*, 1995.
- [4] J. D. Lee and K. A. See, "Trust in automation: Designing for appropriate reliance," *Human Factors*, 2004.
- [5] D. J. Solove, "A taxonomy of privacy," *Univ. of Pennsylvania Law Review*, vol. 154, no. 3, pp. 477–570, 2006.
- [6] A. Cavoukian, "Privacy by Design: The 7 Foundational Principles," *Information and Privacy Commissioner of Ontario*, 2009.
- [7] H. Nissenbaum, *Privacy in Context: Technology, Policy, and the Integrity of Social Life*. Stanford Univ. Press, 2009.
- [8] S. Slade and P. Prinsloo, "Learning analytics: Ethical issues and dilemmas," *American Behavioral Scientist*, vol. 57, no. 10, pp. 1510–1529, 2013.
- [9] K. Crawford and J. Schultz, "Big data and due process: Toward a framework to redress predictive privacy harms," *Boston College Law Review*, 2014.

- [10] F. Pasquale, *The Black Box Society: The Secret Algorithms That Control Money and Information*. Harvard University Press, 2015.
- [11] M. Hildebrandt, "Smart technologies and the end(s) of law," Edward Elgar Publishing, 2015.
- [12] K. O'Neil, *Weapons of Math Destruction*. Crown, 2016.
- [13] European Parliament and Council of the European Union, "Regulation (EU) 2016/679 (General Data Protection Regulation)," *Official Journal of the European Union*, 2016.
- [14] J. Penney, "Chilling effects: Online surveillance and Wikipedia use," *Berkeley Tech. Law Journal*, vol. 31, no. 1, 2016.
- [15] B. Williamson, "Big Data in Education: The Digital Future of Learning, Policy and Practice," Sage, 2017.
- [16] J. Cheney-Lippold, *We Are Data: Algorithms and the Making of Our Digital Selves*. NYU Press, 2017.
- [17] S. Kokolakis, "Privacy paradox heuristics and strategies: A review of references," *Computers & Security*, vol. 64, pp. 122–134, 2017.
- [18] S. U. Noble, *Algorithms of Oppression: How Search Engines Reinforce Racism*. NYU Press, 2018.
- [19] D. Beer, "The data gaze: Capitalism, power and perception," Sage, 2018.
- [20] V. Eubanks, *Automating Inequality*. St. Martin's Press, 2018.
- [21] S. Zuboff, *The Age of Surveillance Capitalism*. PublicAffairs, 2019.
- [22] R. Benjamin, *Race After Technology: Abolitionist Tools for the New Jim Code*. Polity, 2019.
- [23] M. Mitchell et al., "Model Cards for Model Reporting," *FAT* '19*, 2019.
- [24] N. Couldry and U. A. Mejias, "Data Colonialism: Rethinking Big Data's Relation to the Contemporary Subject," *Television & New Media*, 2019.
- [25] S. Andrejevic, "Automated Media," Routledge, 2019.
- [26] L. Amoore, "Cloud Ethics: Algorithms and the Attributes of Ourselves and Others," *Duke University Press*, 2020.