

Architectural Irreversibility: Enforcing Privacy, Governance, and Trust in Intelligent Campus Systems

Dr. S. Suresh Kumar
Principal

Swarnandhra College of Engineering & Technology (Autonomous)
Seetharampuram, Narsapur, Andhra Pradesh, India

Abstract—Intelligent campus infrastructure is growing increasingly reliant on sensors and analytics both to aid in safety and coordination, and to interpret the information flows better. These mechanisms often seek to balance the analytical benefits of observation with institutional obligations to protect privacy. Administrative measures are often in place to prevent abuse and ensure a baseline level of security, but rarely go beyond the capacity to collect, store, or recreate full biodata sets.

This paper introduces the design principle of Architectural Irreversibility: embedding privacy constraints in computations such that reconstruction of the original data is made impossible. By limiting the ability to retrace calculations, Architectural Irreversibility creates new opportunities for enforcing privacy by design. It is suggested here that technical privacy mechanisms that appear invisible to the user rarely serve to build trust, and that interface abstractions must instead be used to communicate privacy constraints directly to stakeholders.

Within this framework, durable public trust emerges from infrastructures that structurally eliminate the ability to observe or retain sensitive information instead of relying solely on administrative restraint.

Index Terms—Architectural Irreversibility, Capability Elimination, Governance-Embedded Inference, Destructive Abstraction, Monotonic Sensitivity Reduction, Socio-Technical Trust.

I. Introduction

The central thesis of this paper is that privacy is best achieved when sensitive observational capability is made structurally impossible rather than simply regulated. Intelligent systems are increasingly prevalent in campus settings, incorporating sensors, analytics, and automated decision-making into everyday functions. While public policy debates rage on, the overall response to pervasive sensing appears to be one of indifference. This, in turn, points to a disconnect between policy initiatives and the public’s underlying unease about potential for surveillance.

A. The Persistence of Reversible Surveillance Architectures

Most contemporary privacy-preserving AI systems encrypt sensitive data but still retain it. They rely on administrative restraint after collection, preserve reconstructable biometrics in encrypted storage, and maintain latent future surveillance capability behind access-control

boundaries. In these architectures, governance is applied to capabilities that have already manifested. The danger is that if such information as biometrics persists in some reconstructable form, its use for other, possibly malevolent purposes will become possible over time. Architectural elimination is different from administrative restriction.

This paper introduces Architectural Irreversibility as a foundational design principle in which strict privacy constraints are embedded directly into the computational data pipeline. Rather than regulating how sensitive data is used after capture, this paradigm structurally constrains whether the system can possess such data at all. Within the ScholarMaster architecture, this subsystem defines irreversible abstraction, capability elimination, governance interposition, and visible boundedness only. It does not validate runtime crash behavior (Paper 18), derive formal adversary theorems (Paper 19), construct distributed compliance logic (Paper 21), synthesize deployment architecture (Paper 20), or perform longitudinal trust analytics (Paper 16).

II. Problem Reframing: Why Privacy-by-Policy Fails

To understand why institutional monitoring systems frequently encounter resistance, it is necessary to examine the structural limitations of policy-driven privacy protection.

A. The False Promise of Post-Hoc Privacy

Many intelligent campus deployments attempt to demonstrate privacy compliance through familiar enterprise security mechanisms. These typically include AES-256 encryption at rest, strict Role-Based Access Control (RBAC), immutable audit logging, and alignment with comprehensive regulatory frameworks like the General Data Protection Regulation (GDPR). Administrators often argue that these mechanisms prevent abuse. However, sociological studies on the use of such systems in educational, medical and smart-city contexts suggest that despite passing various technical security audits, such systems frequently generate distrust and resistance amongst the users [16].

The fundamental problem with this approach is that it is entirely structural; that is, these approaches only come

into play after the sensitive information has been captured, stored, or processed in some way. Encryption protects information that has already been encoded on a hard drive. Access controls dictate who gets to view information that has already been collected. Audit trails provide a forensic record of the actions taken by administrators who have accessed the information. Each of these mechanisms works at the level of the information already present in the system but does nothing to prevent the creation of detailed, high-resolution models of individuals within the system.

From a systems engineering point of view this is potentially dangerous since capability often outweighs governability. Once deployed, and the edge nodes or central servers can sense, store and even infer sensitive data, administrative policies will likely always be playing catch up, trying to stop this capability from being abused. Real-world experience from large-scale technical systems shows that these kinds of safety margins are rarely effective [12]. Normal human error, minor software glitches, external state-sponsored attackers or even just changes in institutional priorities could lead to the initially benign monitoring capabilities being twisted to purposes beyond the original design. While formal policies might restrict what authorized personnel can do with the data, they are unlikely to prevent the data from being collected in the first place.

B. Data as a High-Liability Asset

Within big data infrastructures, raw sensory information often plays a role less akin to an asset and more like a liability to institutional operations. Within the domain of biometric and behavioral analytics, this consideration may be further complicated by the possibility that raw data may take on the characteristics of a high-liability or even hazardous informational asset [17].

If an educational institution is ever faced with storing large volumes of video streams, or raw sound recordings, containing personal details of any of its students, it takes on an enormous responsibility for the protection of that information. Such a database is immediately a tempting target for any potential hackers on the outside, a vital source of information for insider attacks, and a major liability for any potential legal proceedings on the grounds of information disclosure. If the raw information is left in such a state, the institution must expend resources on keeping the data safe from an ever-growing variety of threats.

Constraint-embedded design circumvents this liability entirely by refusing to generate the high-liability asset in the first place. If an edge computing node processes sensory input into irreversible, sparse abstractions and immediately discards the raw photographic captures, a total physical or network compromise of that node yields only coarse semantic signals. Attackers retrieve meaningless

coordinates and timestamps rather than reconstructable facial biometrics, neutralizing the incentive for the attack.

C. Capability Leakage and Mission Creep

Centralized data architectures often accumulate latent analytical capability over time. Even if raw biometric data remains unused initially, the mere existence of archived captures enables vast future inference capabilities. Machine learning models can be effortlessly retrained using historical footage to extract entirely new features that were not considered during the initial deployment. Deep behavioral correlations can be discovered years after the data was originally collected. Feature expansions driven by new administrative demands may introduce entirely new analytical objectives that students never explicitly consented to [21].

This gradual transformation rarely takes place at once, through one single malicious step; more often than not, it is the result of several small tweaks, each of which seems reasonable enough on its own. For instance, one might begin with a module that tracks the library’s usage density to calculate and adjust the heating and air conditioning, and then another one might be added to track how attentive students are during exams. Overall, this process can turn a seemingly innocuous learning environment into a walled garden of surveillance [22]. This growing concern often has little to do with abstract worries about technology or Big Brother-style machinations. On the contrary, it frequently manifests as an intuitively justified suspicion towards procedures that employ asymmetric observation capabilities for essentially paper-and-pencil processes.

D. The Limits of Compliance-Driven Design

Legal and regulatory requirements are often presented as guarantees of safety, but regulations control uses, not existence. They define what is allowed, not what must be done. They specify terms of use, such as the maximum period of retention, but not technical realizations that would make compliance inevitable.

A surveillance system can be fully legal under current law and still capture biometric video data at maximum resolution, as long as such data is secured and purged at the end of the academic year. From a software design perspective, this is a precarious position to be in. The system is vulnerable not only to changes in policy but to the slow creep of time itself. Laws evolve, often in unpredictable ways, meaning that the most optimal solution today may well prove illegal next week. The only way to insulate a program like this from the capriciousness of lawmakers is to make its compliance with privacy laws an architectural guarantee.

III. A Taxonomy of Privacy Architectures

Clarifying the distinction between governance mechanisms and structural constraints requires an initial catego-

rization of privacy architectures. Let us denote the surveillance capability of the system with C , the administrative governance controls with G , and the underlying physical and software architecture with A . The three classes can then be defined as follows:

1. Policy-Constrained Systems: In this case, the system was designed in such a way that it always possesses the full capability (C). It can monitor, collect, and store all the data needed by the users. The governing part of the system (G) only allows the use of this capability through the official documents, terms of service, and other legal means. The system may violate privacy at any moment but is programmed not to do so.

2. Cryptographically-Constrained Systems: The surveillance capability (C) still exists in full fidelity, but G , combined with heavy cryptographic encryption, restricts access to the raw data, allowing only authorized key-holders to view it. While the data is considerably safer from outside attackers or casual insider snooping, the institution itself still retains the ultimate capability to decrypt, analyze, and repurpose the data at its own discretion. In practical terms the capability therefore remains latent within the system.

3. Architecturally-Constrained Systems: In this final paradigm, the capability to perform retrospective surveillance or reconstruct biometric identity simply does not exist within the system state space ($C \notin S$). The architecture (A) itself makes these operations mathematically and physically impossible. The system cannot be steered to disclosure of information, even by a rogue administrator or foreign government with a backdoor capability, since such mathematical ability has been removed from the edge.

This taxonomy highlights the profound difference between merely limiting access to a dangerous capability and eliminating that capability entirely from the design space.

IV. Architectural Irreversibility

A. From Reversible Pipelines to Irreversible Systems

The concept of Architectural Irreversibility suggests a way to make intelligent systems safer by changing how they process sensitive information. Most AI systems today operate as reversible processes. The raw input data is encoded, sometimes stored in a database, and then repeatedly decoded and re-encoded for downstream processing through various transformations inside the network. Even when the system uses different representations for different downstream tasks, these often persist in caches or are retained for later use. This enables easier debugging since you can trace any given piece of data back through the pipeline to see what went wrong if there is an error. It makes it much easier to roll out new features later, as well, because you can always reverse transform any existing data if you need to change how it was initially processed. But it creates unacceptable risks for privacy, security, and safety.

Architecture of Irreversibility, on the other hand, suggests an entirely new approach, one based on transformations that result in destruction. Having transformed at a particular edge node, the system cannot reconstruct the initial sensations from which it was built.




Fig. 1. The Irreversible Pipeline. Raw data exists exclusively in volatile memory buffers that are never committed to persistent storage. The data is destructively transformed at the abstraction boundary, mathematically preventing reconstruction before higher-order inference occurs.

B. The Capability Elimination Principle

The proposed architecture relies on a central engineering principle termed Capability Elimination—the deliberate removal of unnecessary observational capacity. We distinguish three capability states:

- C_{raw} : Raw observational capability exists (high-fidelity biometric data is present and accessible).
- $C_{abstract}$: Data has been transformed but auxiliary information may still permit partial reconstruction.
- $C_{nonreconstructable}$: The transformation is irreversible and all auxiliary side-channel information has been destroyed.

The architectural goal is to enforce the irreversible state transition $C_{raw} \rightarrow C_{abstract} \rightarrow C_{nonreconstructable}$ at the edge, before any persistent storage or network transmission occurs. Consider a mathematical transformation applied to a video frame: $f : X \rightarrow Y$. If f is non-invertible, and the original raw state X is actively destroyed, then reconstruction requires external auxiliary information that no longer exists anywhere in the system.

Attaining the $C_{nonreconstructable}$ state necessitates the simultaneous presence of multiple rigorous engineering controls, including particularly non-invertible transformation functions, the complete eradication of any form of raw persistence, the absence of key-based recovery mechanisms, and the elimination of any reversible dense embeddings. The property of irreversibility is not only algorithmic but also depends on the system’s configuration assumptions, such as secure boot mechanisms and firmware hardening to resist undocumented side-channel persistences.

C. Relationship to Privacy-Enhancing Technologies

It is essential to distinguish between the described architectural approach to privacy protection and the modern Privacy-Enhancing Technologies (PETs). The latter refers to methods like Differential Privacy, Federated Learning, and Secure Hardware Enclaves, which allow for excellent protection of already existing datasets. However, such solutions do not prevent the appearance of raw data in the system because they only protect the data that already exists within a certain boundary. On the contrary, Architectural Irreversibility aims to ensure that specific representations of sensitive information do not appear at all, thus preventing any possible honeypot from being created.

D. The Illusion of De-Identification

Traditional approaches to privacy typically involve some form of anonymization or pseudonymization of private data. Privacy researchers have shown that anonymized data can often be de-anonymized by applying auxiliary information [7]. Pseudonymization also fails to provide strong privacy guarantees since the mapping from pseudonyms to real-world identifiers can be reverse-engineered by anyone with access to it [10]. Even when data is encrypted for release to outside parties, it remains fully useable (in terms of biometric content) to insiders within the institution performing the encryption. Architectural elimination, on the other hand, removes the ability to reconstruct private information from the system environment altogether.

E. Dimensionality Reduction as a Mechanism of Privacy

Within the proposed framework, privacy is embedded into the architecture through aggressively entropic dimensionality reduction. High-resolution media (e.g., 4K video frames containing distinguishable faces, clothing logos, or other characteristic visual attributes) are irreversibly compressed into extremely sparse semantic vectors during the encoding phase.

One example of this is skeletal pose reconstruction. Rather than keeping images of students in the lecture hall, the system stores only the information about the position of each joint, which is then stored as coordinates in relation to one another. In doing so, it both removes any identifying information and keeps the information that the system needs to analyze behavior.

V. Mandatory Governance

A. Governance as Computational Interposition

Effective AI governance requires more than just policy documentation. In many cases, governance takes place elsewhere, in other systems, such as policy management or compliance. The above architecture reflects governance through execution as compiled code. Governance is performed by being interposed computationally in the decision-making chain. Note that in this architecture,

the inference engines are not direct actors. Effective governance would seek to narrow the possible outputs by the model in accordance with institutional policy. By embedding policy in code boundaries, governance is realized in constraints on what outputs can be passed beyond the edge.

B. Epistemic Humility

The machine learning model provides us with probabilistic conclusions and not with ground truth data. Letting the model take any administrative action could mean letting it label a student as disengaged in an attempt to correct a statistical mean [1].

The design of the architecture embodies the principles of Epistemic Humility [8] by separating judgment (the model's statistical inference) from authority (approval to act on) – indeed, the inference engine may be right, but it can only express its judgements internally, not as actions on the World. Only a separate Governance layer, endowed with domain specific institutional policy, can take responsibility for externalising results on the wider institutional environment, if it chooses to do so.

C. Constraint-First Systems Engineering

Traditional AI systems are optimized: they are designed to maximize some objective function. Architectural Irreversibility embraces a constraint-first systems engineering philosophy in which maximum utility is deliberately avoided. The utility function of the system is capped to realize Nissenbaum's contextual integrity [12]. If the policy verification server becomes unreachable (or if the audit database becomes unreachable), the system reverts to a default failed state. Inference results are not returned and execution is halted. The design incorporates limited observability, irreversible reduction, fail-closed operation, and constrained inference authorities.

VI. Visible Privacy and Trust

A. Transparency vs Legibility

Even mathematically rigorous privacy assurances might fail to provide the illusion of control that make them appealing to laypeople if they are opaque to system participants. In the development of socio-technical systems, the values of transparency and legibility are frequently conflated, with potentially hazardous results. A public repository containing the source code for a system may well be transparent in a technical sense, but it does nothing to reassure a non-technical student that she is not being monitored in a classroom where monitoring takes place. Students who are subjected to surveillance infrastructure are unlikely to have the time or technical expertise to review relevant source code repositories – their beliefs about the nature of the system must instead be formed on the basis of their direct observations of its local effects [4].

B. The Semiotics of Abstraction

To this end, the system design imposes significant restrictions on the user-facing interfaces, constraining them to abstract representations. Administrative interfaces do not provide security personnel with live video feed previews but rather utilize skeletons, nodes, or other graphical indicators placed on an empty grid [2].

This visual language communicates an incredibly important, immediate constraint: that identity information is physically inaccessible to the human operators facing the screen. The visual has dual value as both a technical privacy safeguard and a trust-building mechanism, short-circuiting potential mistrust by enabling non-technical users to assess the situation for themselves. This avoids the situation where privacy is only maintained through the unilateral assurance of one party and the unknowledgeable acceptance of another.

VII. The Architecture of Stewardship

The principles established above (irreversibility of transformations, mandatory governance, and visibility of privacy) can be implemented through a progressive reduction of identity dimensions in a layered fashion as data is processed up the stack. The detailed design of such a system, including derivation of formal layer specifications, execution contracts between layers, and deployment topologies, is the subject of companion work on reference architecture design (Paper 20). This paper establishes the doctrinal foundation that informs that architectural specification and its constraints.

A. Monotonic Sensitivity Reduction as Governing Invariant

From an abstract perspective, the pipeline performs a monotonic reduction in confidence. This is the system's basic principle, its *raison d'être*. The environment provides input data of a certain quality, from which the system must extract specific values. Further processing involves a progressive loss of identity characteristics. It is essential to note that these characteristics cannot be recovered at later stages. In other words, each subsequent level of abstraction possesses a lower information content compared to the previous one. Ultimately, the information reaches an abstract level, where it loses all individual characteristics and becomes a set of anonymous signals.

The key doctrinal constraints governing this pipeline are:

- **Monotonic Sensitivity Reduction:** Information sensitivity must decrease at every boundary transition. No layer may re-enrich data with identity dimensions that a prior boundary destroyed. This property must hold consistently across pipeline discussion, governance reasoning, and systems framing.
- **Unidirectional Flow:** Only data originating within the system can proceed up the pipeline. Feedback loop data cannot pass back down the pipeline from a higher

layer if it contains information that could be used to enhance the analysis of lower-layer captured data.

- **Mandatory Governance Interposition:** No abstracted output may reach external sinks without passing through an explicit, computationally interposed governance evaluation.
- **Fail-Closed Default:** Any boundary failure causes the pipeline to halt rather than degrade into a permissive state.

VIII. Non-Negotiable Architectural Constraints

Traditional software systems are controlled by a highly malleable set of legal contracts and Terms of Service. This thesis suggests designing systems with non-negotiable architectural invariants that forbid what the system may do, no matter who is asking or what their level of administrative privilege is.

Boundary Irreversibility is a principle that ensures that raw sensory data cannot be communicated beyond the abstraction boundary into storage. Governance Liveness is the principle that ensures that nothing can be generated or communicated as an output without undergoing active policy consideration and approval. Algorithmic Forgetting is the principle that ensures that context is subject to a rolling expiration so that it cannot be used to build up lasting psychological profiles of participants [14]. Finally, the Principle of Least Privilege in Memory ensures that the only memory with any privileged access to these perception layers is strictly limited. In these terms, the described limitations are precisely the mechanisms by which the author intended for the machine to remain within its boundaries.

IX. Ethical Positioning

A. From Compliance to Constraint

Architecture irreversibility is a shift from the policy of control to the policy of constraints to system design [17]. The control policy allows systems to do what they want but not do what is necessary. Architectural constraints set boundaries about what a system must not be able to do. By building data minimization into the design of a system's software and hardware, we can reduce society's reliance on the sometimes-erroneous, sometimes-malicious choices of its administrators [8].

B. The Stewardship Model

Automated Stewardship differs profoundly from many surveillance-like information gathering approaches, where the surveillance architectures are geared towards gaining maximum information retention, through systematic obfuscation of information, in order to enforce a certain power differential over the observed objects [19]. The stewardship paradigm, on the other hand is aimed at reducing potential for harm, and gaining acceptance through visibility and constraint [13].

C. Practical Trade-offs and Operational Realism

Irreversible systems pose several challenges concerning their practical use. The debugging of such systems becomes significantly more complicated due to the absence of raw data for analysis in case of failure. Besides, some research activities that require analyzing past events using biometrics data become impossible. Notably, these restrictions are imposed to ensure privacy, without which biometrics would have been even more intrusive. Finally, the management of such systems can be complicated by the necessity to maintain the integrity of devices, trust their firmware, and standardize procedures across different hardware platforms at the university level.

D. Limitations

The architectural enforcement of irreversibility and governance constraint introduces several operational boundaries. Under the defined architectural boundary, the following constraints apply:

- **Diagnostic Opacity:** By destructively transforming raw video inputs at the edge, forensic debugging of model misclassifications becomes nearly impossible, as the ground-truth visual data no longer exists.
- **Longitudinal Rigidity:** Irreversibility fundamentally prevents retrospective feature expansion; institutions cannot retroactively train new machine learning models on historical environmental conditions because the raw data was never retained.
- **Hardware Trust Dependency:** The guarantee of irreversibility is based on the physical security of the edge sensing nodes; compromised firmware could potentially bypass the abstraction barrier before destruction takes place. Side-channel leakage is still possible within the constrained execution model.
- **Side-Channel Vulnerability:** Even in the case of raw biometrics, using more fine-grained skeletal vectors or metadata could still carry risks of exposing behavioral patterns that may be exploited in complex ways to gain insight into one's personal identifiers, despite the reduced reconstructability of the information.
- **Asymmetric Development Costs:** Embedding governance directly into compiled pipeline boundaries significantly increases engineering complexity and maintenance cost compared to traditional, permissive data-lake architectures.

X. Conclusion

This paper presented Architectural Irreversibility as a bounded architectural doctrine for structural capability elimination in privacy-sensitive campus infrastructures. By embedding one-way abstraction boundaries directly into the computational pipeline, the architecture reduces latent surveillance capability rather than administratively restricting its exercise. The irreversible state transition $C_{raw} \rightarrow C_{abstract} \rightarrow C_{nonreconstructable}$, governed by monotonic sensitivity reduction, ensures that information

sensitivity decreases at every pipeline boundary under the defined architectural assumptions.

Visible abstraction interfaces reduce informational asymmetry by allowing non-technical users to observe the system's operational limits directly. Together these mechanisms demonstrate how privacy can be treated as a foundational property of system architecture rather than a retrospective regulatory obligation.

Within the ScholarMaster architecture, this framework operates exclusively as the architectural governance and irreversible privacy subsystem—defining capability elimination, governance interposition, and visible boundedness without validating runtime crash behavior (Paper 18), deriving formal adversary theorems (Paper 19), constructing distributed compliance logic (Paper 21), synthesizing deployment topology (Paper 20), or performing longitudinal trust analytics (Paper 16).

References

- [1] J. Sweller, "Cognitive load during problem solving," *Cognitive Science*, vol. 12, no. 2, pp. 257–285, 1988.
- [2] P. Milgram and F. Kishino, "A taxonomy of mixed reality displays," *IEICE Trans. Inf. Syst.*, vol. E77-D, no. 12, pp. 1321–1329, 1994.
- [3] J. Bentham, *The Panopticon Writings*. Verso, 1995.
- [4] M. R. Endsley, "Toward a theory of situation awareness," *Human Factors*, vol. 37, no. 1, pp. 32–64, 1995.
- [5] R. C. Mayer, J. H. Davis, and F. D. Schoorman, "An integrative model of organizational trust," *Acad. of Management Review*, vol. 20, no. 3, pp. 709–734, 1995.
- [6] R. T. Azuma, "A survey of augmented reality," *Presence*, vol. 6, no. 4, pp. 355–385, 1997.
- [7] L. Sweeney, "k-anonymity: A model for protecting privacy," *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*, vol. 10, no. 5, pp. 557–570, 2002.
- [8] S. Jasanoff, "Technologies of Humility: Citizen Participation in Governing Science," *Minerva*, vol. 41, no. 3, pp. 223–244, 2003.
- [9] J. D. Lee and K. A. See, "Trust in automation: Designing for appropriate reliance," *Human Factors*, vol. 46, no. 1, pp. 50–80, 2004.
- [10] D. J. Solove, "A taxonomy of privacy," *Univ. of Pennsylvania Law Review*, vol. 154, no. 3, pp. 477–570, 2006.
- [11] A. Cavoukian, "Privacy by design: The 7 foundational principles," *Information and Privacy Commissioner of Ontario*, 2009.
- [12] H. Nissenbaum, *Privacy in Context: Technology, Policy, and the Integrity of Social Life*. Stanford Univ. Press, 2009.
- [13] L. Floridi, *The Ethics of Information*. Oxford University Press, 2013.
- [14] R. R. Hoffman et al., "Accelerated proficiency and facilitated retention," *Cognitive Processing*, vol. 14, no. 2, pp. 147–163, 2013.
- [15] M. Hildebrandt, *Smart Technologies and the End(s) of Law*. Edward Elgar, 2015.
- [16] F. Pasquale, *The Black Box Society*. Harvard Univ. Press, 2015.
- [17] B. Schneier, *Data and Goliath: The Hidden Battles to Collect Your Data and Control Your World*. W. W. Norton & Company, 2015.
- [18] J. Penney, "Chilling effects: Online surveillance and Wikipedia use," *Berkeley Tech. Law Journal*, vol. 31, no. 1, 2016.
- [19] S. U. Noble, *Algorithms of Oppression*. NYU Press, 2018.
- [20] V. Eubanks, *Automating Inequality*. St. Martin's Press, 2018.
- [21] N. Couldry and U. A. Mejias, "Data colonialism," *Television & New Media*, vol. 20, no. 4, pp. 336–349, 2019.
- [22] S. Zuboff, *The Age of Surveillance Capitalism*. PublicAffairs, 2019.