

Generalized Cross-Modal Recovery under Compromised Primary Sensing

Dr. S. Suresh Kumar

Principal, Swarnandhra College of Engineering & Technology (Autonomous)
Seetharampuram, Narsapur, Andhra Pradesh, India — principal@swarnandhra.ac.in

Abstract—Cyber-physical edge systems operating in unconstrained environments frequently encounter physical sensor degradation, optical occlusions, and adversarial tampering. While multi-modal sensory architectures incorporate complementary sensing streams (e.g., optical RGB, skeletal pose kinematics, and acoustic spectral features), conventional multi-modal fusion mechanisms employ static weighting matrices that catastrophically propagate noise from a compromised primary channel into the joint feature representation. In this paper, we present the information-theoretic formulation, mathematical analysis, and empirical verification of Layer-1 Generalized Cross-Modal Consensus Recovery for edge intelligence. We introduce a symmetric Jensen-Shannon Divergence (JSD) formulation that continuously measures inter-modality distribution discrepancy against a dynamic arithmetic mixture consensus distribution. We prove from first principles that the symmetric JSD is strictly bounded in $[0, \ln 2]$, derive Pinsker-type total variation inequality bounds $(\frac{1}{2}\|P - Q\|_{TV})^2 \leq \text{JSD}(P \parallel Q) \leq \ln(2)\|P - Q\|_{TV}$, and analyze infinitesimal Fisher information metric geometry on the statistical probability simplex ($ds_{FR}^2 = 8 \cdot \text{JSD} + O(\|dP\|^3)$). To resolve temporal latency mismatches across heterogeneous hardware sensors, we formalize an asynchronous multi-rate ring buffer synchronization mechanism aligning 30 FPS RGB video, 100 Hz IMU kinematics, and 15 FPS acoustic spectral frames with software phase-locked loop (PLL) timestamp compensation. Evaluated across progressive visual degradation regimes (0%, 20%, 50%, and 80% synthetic sensory corruption), our dynamic consensus mechanism achieves a 100% (1.0000) state recovery rate under single-channel degradation. Under severe 80% visual corruption, where single-modality RGB accuracy collapses from 1.0000 to 0.1867, the dynamic trust weight of the corrupted optical channel automatically decays from 0.4000 to 0.0500, autonomously transferring decision authority to intact acoustic and skeletal channels (0.4750 each). We formalize multi-channel breakdown boundaries and prove that dynamic consensus guarantees continuous state survivability under single-modality sensory compromise.

Keywords—Cross-modal recovery, sensor fusion, Jensen-Shannon divergence, dynamic trust weighting, multi-rate synchronization, Fisher information metric, edge survivability.

\begin{abstract}

Cyber-physical edge systems operating in unconstrained environments frequently encounter physical sensor degradation, optical occlusions, and adversarial tampering. While multi-modal sensory architectures incorporate complementary sensing streams (e.g., optical RGB, skeletal pose kinematics, and acoustic spectral

features), conventional multi-modal fusion mechanisms employ static weighting matrices that catastrophically propagate noise from a compromised primary channel into the joint feature representation. In this paper, we present the information-theoretic formulation, mathematical analysis, and empirical verification of Layer-1 Generalized Cross-Modal Consensus Recovery for edge intelligence. We introduce a symmetric Jensen-Shannon Divergence (JSD) formulation that continuously measures inter-modality distribution discrepancy against a dynamic arithmetic mixture consensus distribution. We prove from first principles that the symmetric JSD is strictly bounded in $[0, \ln 2]$, derive Pinsker-type total variation inequality bounds $(\frac{1}{2}\|P - Q\|_{TV})^2 \leq \text{JSD}(P \parallel Q) \leq \ln(2)\|P - Q\|_{TV}$, and analyze infinitesimal Fisher information metric geometry on the statistical probability simplex ($ds_{FR}^2 = 8 \cdot \text{JSD} + O(\|dP\|^3)$). To resolve temporal latency mismatches across heterogeneous hardware sensors, we formalize an asynchronous multi-rate ring buffer synchronization mechanism aligning 30 FPS RGB video, 100 Hz IMU kinematics, and 15 FPS acoustic spectral frames with software phase-locked loop (PLL) timestamp compensation. Evaluated across progressive visual degradation regimes (0%, 20%, 50%, and 80% synthetic sensory corruption), our dynamic consensus mechanism achieves a 100% (1.0000) state recovery rate under single-channel degradation. Under severe 80% visual corruption, where single-modality RGB accuracy collapses from 1.0000 to 0.1867, the dynamic trust weight of the corrupted optical channel automatically decays from 0.4000 to 0.0500, autonomously transferring decision authority to intact acoustic and skeletal channels (0.4750 each). We formalize multi-channel breakdown boundaries and prove that dynamic consensus guarantees continuous state survivability under single-modality sensory compromise.

\end{abstract}

\begin{IEEEkeywords}

Cross-modal recovery, sensor fusion, Jensen-Shannon divergence, dynamic trust weighting, multi-rate synchronization, Fisher information metric, edge survivability.

\end{IEEEkeywords}

INTRODUCTION

Autonomous cyber-physical systems, such as automated transit hubs, industrial robotic inspection units, autonomous surveillance arrays, and smart campus access control nodes, depend on continuous, high-fidelity state estimation to enforce real-time security

and compliance [Ref]. However, single-sensor deployments are inherently brittle: primary optical cameras are acutely vulnerable to lens defocus blur, illumination extremes, atmospheric fog, lens smearing, physical occlusions, and adversarial spoofing probes [Ref].

To mitigate single-channel vulnerability, modern edge platforms deploy heterogeneous sensor suites combining optical RGB cameras, skeletal pose estimators, thermal infrared detectors, acoustic arrays, and inertial measurement units (IMUs) [Ref]. Nevertheless, simply aggregating multi-modal inputs does not guarantee system robustness. Standard multimodal architectures suffer from severe architectural limitations:

- **Early Feature Concatenation:** Concatenating raw sensory embeddings into a unified dense layer creates tight feature entanglement. When an optical channel degrades, corrupted feature values propagate through subsequent layers, corrupting the entire latent representation [Ref].
- **Static Late Fusion:** Linearly combining independent decision vectors via fixed heuristic weights (e.g., $w_{\text{rgb}}=0.5$, $w_{\text{audio}}=0.5$) implicitly assumes time-invariant sensor reliability. Under severe optical noise, static weighting forces the corrupted visual prediction to pull down the joint decision, causing system accuracy to collapse [Ref].
- **Deep Attention Overhead:** Heavy transformer-based cross-attention mechanisms (e.g., Multimodal Transformers [Ref]) model all-to-all token interactions but introduce prohibitive computational latencies (>40 ms) and uncalibrated attention logits under out-of-distribution noise [Ref].

To overcome these deficiencies, we draw a fundamental distinction between *Multimodal Fusion* (aggregating features assuming clean inputs) and *Multimodal Recovery* (actively identifying corrupted channels, suppressing their authority, and reconstructing intact state estimations without model retraining). The central scientific question addressed in this paper is:

$\begin{quote}$

When a primary sensory channel experiences severe degradation, how can an edge system detect inter-modality disagreement in real time, dynamically decay the compromised channel's authority, and preserve accurate downstream state estimation using intact secondary modalities?

$\end{quote}$

This paper establishes the mathematical foundations, synchronization architecture, and empirical verification of *Generalized Cross-Modal Consensus Recovery*. Our architecture treats modality trust as a continuous, time-varying dynamic state governed by symmetric Jensen-Shannon Divergence (JSD). When sensory noise or physical occlusion corrupts a sensor, its output distribution diverges from the arithmetic consensus distribution P_c , triggering exponential trust weight decay and autonomously transferring decision authority to intact secondary streams.

Specifically, this paper makes four core contributions:

- **Information-Theoretic Consensus & Boundedness Proof:** We formalize cross-modal agreement on the statistical probability simplex using symmetric Jensen-Shannon Divergence. We prove from first principles that $JSD(P_m \parallel P_c)$ is strictly bounded in $[0, \ln 2]$, derive Pinsker-type total variation bounds, and prove local equivalence to the Riemannian Fisher-Rao metric tensor.
- **Dynamic Trust Gradient Adaptation:** We formulate an exponential trust weighting function with sensitivity parameter $\beta=5.0$, deriving the analytical self-gradients ($\frac{\partial}{\partial w_m} JSD_m < 0$) and cross-gradients ($\frac{\partial}{\partial w_m} JSD_j > 0$) that govern smooth, monotonic authority reallocation.
- **Asynchronous Multi-Rate Synchronization Architecture:** We formulate a multi-rate ring buffer reference model with software phase-locked loop (PLL) clock tracking to align heterogeneous sensor streams (30 FPS RGB, 100 Hz IMU, 15 FPS Audio), distinguishing the theoretical PLL reference model from the production runtime's 1.0 s timestamp skew checking window in

ConsistencyChecker.

- **Empirical Recovery & Breakdown Boundaries:** Across progressive synthetic degradation levels (0%, 20%, 50%, 80%), we demonstrate that dynamic consensus achieves a 100% (1.0000) state recovery rate under single-channel degradation, dynamically attenuating corrupted RGB trust from 0.4000 $\rightarrow$ 0.0500 while intact secondary modalities increase to 0.4750 each. We derive analytical multi-channel breakdown boundaries and fail-closed quarantine criteria.

RELATED WORK & MULTI-MODAL FUSION TAXONOMY

Classical Multisensor Data Fusion

Multisensor data fusion has a long history in distributed radar networks, robotic odometry, and navigation systems [Ref]. Classical methods, such as Extended Kalman Filters (EKF), Unscented Kalman Filters (UKF), and covariance intersection [Ref], fuse continuous state estimates by weighting sensory observations inversely proportional to their covariance matrices. However, classical Kalman fusion assumes zero-mean Gaussian measurement noise and linear/linearized observation models. In modern edge vision where sensor corruption manifests as complex, high-dimensional non-Gaussian feature distortions, covariance-based weighting fails to detect semantic distribution shift.

Deep Multimodal Fusion Strategies

Modern deep multimodal architectures are traditionally categorized by their fusion stage [Ref]:

- **Early Fusion:** Concatenates raw feature maps or sensor tokens at the network input layer [Ref]. Early fusion enables rich low-level cross-modal feature correlations but exhibits zero fault tolerance: a single noisy sensor contaminates the shared feedforward pipeline.
- **Late Fusion:** Trains separate unimodal neural networks and aggregates their softmax outputs via fixed linear averaging or voting [Ref]. While modular, static late fusion cannot dynamically adjust its weights when environmental conditions degrade a specific sensor.
- **Intermediate / Cross-Attention Fusion:** Employs pairwise self-attention or cross-modal transformer blocks (e.g., Multimodal Transformer [Ref] or Perceiver [Ref]) to dynamically weight cross-modal tokens. However, transformer attention matrices scale

quadratically with sequence length, generating execution latencies (>40 ms) that violate edge SLA constraints. Furthermore, attention weights over unbounded logits can allocate high attention to high-frequency noise [Ref].

\begin{table*}[t]

\centering

\caption{Comparative Taxonomy of Multimodal Sensor Fusion, Missing-Modality Learning, and Recovery Paradigms}

\label{tab:fusion_taxonomy}

\resizebox{\textwidth}{!}{\%

\begin{tabular}{@{}lcccc@{}}

→prule

Fusion Paradigm & Weighting Mechanism & Mathematical Bound & Degradation Resilience & Multi-Rate Sync & Edge Overhead & Exact Limitation / P24 Differentiator \midrule

1. Early Fusion (Concat) [Ref] & Fixed Uniform & None & Catastrophic Collapse & No & Minimal (<0.1 ms) & Noise propagates joint representation \
2. Late Fusion (Averaging) [Ref] & Fixed Linear Weights & None & Linear Degradation & Partial & Low (0.5 ms) & Static weights cannot isolate failing sensors \
3. Cross-Modal Transformer [Ref] & Self-Attention Matrix & Unbounded Logits & Moderate (Attention) & No & High (>40 ms) & Quadratic complexity, uncalibrated attention logits \
4. Generative Imputation [Ref] & Latent Reconstruction & Reconstruction Loss & Moderate (Hallucination) & No & Heavy (15–25 ms) & Hallucinates synthetic features under severe noise \
5. Modality Dropout [Ref] & Stochastic Zero-Masking & None & Moderate (65%–75%) & No & Low (0.6 ms) & Static robustness; no runtime trust redistribution \
6. Reliability-Gated Fusion [Ref] & Heuristic Softmax & Uncalibrated & Moderate (70%–80%) & No & Low (0.8 ms) & Heuristic weights lack bounded divergence metric \
7. ScholarMaster JSD Consensus (Ours) & Dynamic JSD Gradient & Strict ($0 \leq \text{JSD} \leq \ln 2$) & Complete (100% Recovery) & Yes (Ring Buffer) & Optimal (1.1 ms) & Bounded divergence, real-time trust adaptation _tomrule

Missing-Modality Learning & Generative Imputation

When a sensor fails or disconnects, missing-modality frameworks attempt to reconstruct the missing channel using generative models, such as variational autoencoders (e.g., SMIL [Ref], Private-Shared VAEs [Ref]) or generative adversarial networks. Although generative imputation restores missing feature dimensions, it incurs significant computational overhead (15–25 ms) and introduces the danger of generative hallucination: synthesizing high-confidence biometric or kinematic features that do not correspond to physical reality. Our consensus recovery framework operates strictly via discriminative trust redistribution without generative hallucinations.

Modality Dropout & Reliability-Weighted Fusion

Modality dropout techniques (e.g., ModDrop [Ref]) stochastically drop input modalities during training to encourage unimodal

independence. While ModDrop improves general network robustness, the deployed inference graph remains static and cannot adjust channel authority based on instantaneous sensory corruption. Reliability-weighted fusion methods [Ref] apply heuristic softmax weighting to model confidence scores. However, because deep models are uncalibrated under out-of-distribution perturbations [Ref], heuristic confidence scores fail to provide bounded, stable authority reallocation.

Information-Theoretic Divergence & Consensus

Information-theoretic metrics, notably Kullback-Leibler (KL) divergence, are widely used to measure statistical distance between probability distributions [Ref]. However, classical KL divergence is asymmetric and diverges to infinity when distribution supports do not strictly overlap. The Jensen-Shannon Divergence (JSD) [Ref] addresses these theoretical shortcomings by averaging relative entropies against a symmetric mixture distribution. Our architecture exploits the mathematical boundedness and metric properties of symmetric JSD on the probability simplex to drive a closed-loop dynamic trust adaptation engine for edge intelligence. Table~1 synthesizes these seven paradigms against our proposed architecture.

INFORMATION-THEORETIC JSD CONSENSUS FORMULATION

Modality Probability Representations & Mixture Consensus

Let $M = \{m_1, m_2, \dots, m_{|M|}\}$ denote the active set of heterogeneous sensory modalities deployed on the edge platform (in our primary realization, $|M|=3$: m_1 = Optical RGB, m_2 = Acoustic, and m_3 = Pose). Each modality $m \in M$ processes incoming sensor frames to output a normalized probability distribution $P_m \in \Delta^K$ over K discrete semantic hypothesis states:

$$\text{Equation: } P_m = [P_m(1), \dots, P_m(K)]^T, \quad \sum_{k=1}^K P_m(k) = 1.$$

To establish an unbiased reference baseline against which individual channel disagreement can be evaluated, we construct the instantaneous arithmetic mixture consensus distribution $P_c \in \Delta^K$:

$$\text{Equation: } P_c(k) = \frac{1}{|M|} \sum_{m \in M} P_m(k), \quad \forall k \in \{1, \dots, K\}.$$

Justification for Arithmetic Mixture An arithmetic mixture $P_c = \frac{1}{|M|} \sum_m P_m$ is strictly preferred over geometric pooling ($P_c \propto \prod_m P_m^{1/|M|}$) because

arithmetic pooling preserves convex linear closure on the simplex Δ^K and ensures non-empty common support. Specifically, if a corrupted sensor erroneously assigns $P_{m_1}(k) = 0$ to the true class k^* , geometric pooling would annihilate the consensus probability ($P_c(k^*) = 0$), whereas arithmetic pooling preserves positive support from intact secondary sensors ($P_c(k^*) \geq \frac{1}{|M|} P_{m_2}(k^*) > 0$).

First-Principles Proof of Symmetric JSD Boundedness

Definition: [Jensen-Shannon Divergence]

Let $P_m, P_c \in \Delta^K$ be discrete probability distributions on the K -dimensional simplex. The Jensen-Shannon Divergence between modality distribution P_m and the consensus distribution P_c is defined as:

$$\text{Equation: } JSD(P_m \parallel P_c) = \frac{1}{2} KL(P_m \parallel \bar{M}_m) + \frac{1}{2} KL(P_c \parallel \bar{M}_m),$$

where $\bar{M}_m = \frac{1}{2}(P_m + P_c)$ is the equal mixture distribution, and $KL(P \parallel Q) = \sum_{k=1}^K P(k) \ln \frac{P(k)}{Q(k)}$ is the Kullback-Leibler divergence with the natural logarithm convention ($\ln$).

End of Definition

Theorem: [JSD Information-Theoretic Bounds]

Label: thm:jsd_bounds

For any discrete probability distributions $P_m, P_c \in \Delta^K$, the Jensen-Shannon Divergence is symmetric, non-negative, and strictly bounded:

$$\text{Equation: } 0 \leq JSD(P_m \parallel P_c) \leq \ln(2) \approx 0.69315 \text{ nats.}$$

End of Theorem

Proof:

We express $JSD(P_m \parallel P_c)$ in terms of the Shannon entropy functional $H(P) = -\sum_{k=1}^K P(k) \ln P(k)$:

Align:

$$JSD(P_m \parallel P_c) = \frac{1}{2} \sum_{k=1}^K P_m(k) \ln \frac{P_m(k)}{\bar{M}_m(k)} + \frac{1}{2} \sum_{k=1}^K P_c(k) \ln \frac{P_c(k)}{\bar{M}_m(k)}$$

$$= \frac{1}{2} \sum_{k=1}^K P_m(k) \ln P_m(k) - \frac{1}{2} \sum_{k=1}^K P_m(k) \ln \bar{M}_m(k) + \frac{1}{2} \sum_{k=1}^K P_c(k) \ln P_c(k) - \frac{1}{2} \sum_{k=1}^K P_c(k) \ln \bar{M}_m(k)$$

$$= -\frac{1}{2} H(P_m) - \frac{1}{2} H(P_c) + \frac{1}{2} \sum_{k=1}^K P_m(k) \ln \bar{M}_m(k) + \frac{1}{2} \sum_{k=1}^K P_c(k) \ln \bar{M}_m(k)$$

$$= -\frac{1}{2} H(P_m) - \frac{1}{2} H(P_c) + \frac{1}{2} \sum_{k=1}^K P_m(k) \ln \bar{M}_m(k) + \frac{1}{2} \sum_{k=1}^K P_c(k) \ln \bar{M}_m(k)$$

$$= -\frac{1}{2} H(P_m) - \frac{1}{2} H(P_c) + \frac{1}{2} \sum_{k=1}^K P_m(k) \ln \bar{M}_m(k) + \frac{1}{2} \sum_{k=1}^K P_c(k) \ln \bar{M}_m(k)$$

End of Align

Lower Bound ($0 \leq JSD$): By the strict concavity of the Shannon entropy functional $H(P)$ on the simplex Δ^K , Jensen's inequality states:

$$\text{Equation: } H\left(\frac{P_m + P_c}{2}\right) \geq \frac{1}{2} H(P_m) + \frac{1}{2} H(P_c).$$

Subtracting $\frac{1}{2} H(P_m) + \frac{1}{2} H(P_c)$ from both sides in (1) yields $JSD(P_m \parallel P_c) \geq 0$. By strict concavity of entropy, equality $JSD = 0$ holds if and only if $P_m = P_c$ almost everywhere.

Upper Bound ($JSD \leq \ln 2$): Expanding the mixture entropy:

Align:

$$H\left(\frac{P_m + P_c}{2}\right) = -\sum_{k=1}^K \frac{P_m(k) + P_c(k)}{2} \ln \frac{P_m(k) + P_c(k)}{2}$$

$$= -\ln(2) - \frac{1}{2} \sum_{k=1}^K P_m(k) \ln(P_m(k) + P_c(k)) - \frac{1}{2} \sum_{k=1}^K P_c(k) \ln(P_m(k) + P_c(k)).$$

$$= -\ln(2) - \frac{1}{2} \sum_{k=1}^K P_m(k) \ln(P_m(k) + P_c(k)) - \frac{1}{2} \sum_{k=1}^K P_c(k) \ln(P_m(k) + P_c(k)).$$

End of Align

Because $P_m(k) + P_c(k) \geq P_m(k)$ and $P_m(k) + P_c(k) \geq P_c(k)$ for all k , the monotonicity of the logarithm implies:

Align:

$$-\ln(P_m(k) + P_c(k)) \leq -\ln P_m(k),$$

$$-\ln(P_m(k) + P_c(k)) \leq -\ln P_c(k).$$

End of Align

Substituting these bounds into the entropy expansion:

Align:

$$H\left(\frac{P_m + P_c}{2}\right) \leq -\ln(2) - \frac{1}{2} \sum_{k=1}^K P_m(k) \ln P_m(k) - \frac{1}{2} \sum_{k=1}^K P_c(k) \ln P_c(k)$$

$$= -\ln(2) - \frac{1}{2} H(P_m) - \frac{1}{2} H(P_c).$$

$$= -\ln(2) - \frac{1}{2} H(P_m) - \frac{1}{2} H(P_c).$$

End of Align

Substituting this upper bound into (1) directly yields:

$$\text{Equation: } JSD(P_m \parallel P_c) \leq \ln(2) \approx 0.69315 \text{ nats.}$$

Equality $JSD = \ln(2)$ is attained if and only if the distributions P_m and P_c have mutually disjoint supports ($P_m \perp P_c$, i.e., $P_m(k) \cdot P_c(k) = 0$ for all k).

End of Proof

Pinsker-Type Total Variation Metric Bounds

To relate the information divergence JSD directly to standard statistical distance metrics, we establish its relationship with Total Variation (TV) distance.

Corollary: [Total Variation Metric Bounds]

Label: cor:tv_bounds

Let $\|P_m - P_c\|_{TV} = \frac{1}{2} \sum_{k=1}^K |P_m(k) - P_c(k)| \in [0, 1]$ denote the total variation distance. Under the natural logarithm convention, the Jensen-Shannon Divergence satisfies the two-sided inequalities:

$$\text{Equation: } \frac{1}{2} \|P_m - P_c\|_{TV}^2 \leq JSD(P_m \parallel P_c) \leq \ln(2) \|P_m - P_c\|_{TV}.$$

End of Corollary

Proof:

By classical Pinsker's inequality applied to Kullback-Leibler divergence [Ref]:

$$\begin{aligned} & \text{KL}(P_m \parallel \bar{M}_m) \geq 2 \|P_m - \bar{M}_m\|_{TV}^2 \\ & \leq 2 \left(\frac{1}{2} \|P_m - P_c\|_{TV} \right)^2 \\ & = \frac{1}{2} \|P_m - P_c\|_{TV}^2. \end{aligned}$$

Summing the symmetrical components yields $\text{JSD}(P_m \parallel P_c) \geq \frac{1}{2} \|P_m - P_c\|_{TV}^2$. The upper bound follows from the convexity of JSD with maximal slope $\ln(2)$ at extreme variation $\|P_m - P_c\|_{TV} = 1$ [Ref].

$\text{end}\{\text{proof}\}$

Infinitesimal Fisher Information Geometry

Let the probability simplex be parameterized by local coordinates $\theta \in \mathbb{R}^{K-1}$. On the statistical Riemannian manifold endowed with the Fisher information metric tensor $g_{ij}(\theta) = \sum_{k=1}^K \frac{1}{P_\theta(k)} \frac{\partial P_\theta(k)}{\partial \theta_i} \frac{\partial P_\theta(k)}{\partial \theta_j}$, the infinitesimal squared geodesic distance $ds_{FR}^2 = \sum_{i,j} g_{ij}(\theta) d\theta_i d\theta_j$ satisfies:

Equation: $ds_{FR}^2 = 8 \cdot \text{JSD}(P_\theta \parallel P_{\theta'}) + O(\|d\theta\|^3)$.

This establishes that for small sensory fluctuations ($d\theta \rightarrow 0$), JSD is proportional to the Riemannian metric tensor on the statistical manifold, providing optimal sensitivity to local distribution drift while maintaining global boundedness in $[0, \ln 2]$.

Dynamic Trust Weight Gradient Adaptation

Using the strictly bounded divergence metric $\text{JSD}_m = \text{JSD}(P_m \parallel P_c)$, we define the normalized exponential dynamic trust weight w_m for each modality $m \in M$:

Equation: $w_m = \frac{\exp(-\beta \cdot \text{JSD}_m)}{\sum_{j \in M} \exp(-\beta \cdot \text{JSD}_j)}$,

where $\beta > 0$ is the sensitivity hyperparameter ($\beta = 5.0$).

$\text{begin}\{\text{proposition}\}$ [Analytical Trust Weight Gradients]

The dynamic trust weight allocation satisfies the following analytical gradient properties:

$\text{begin}\{\text{align}\}$

Self-Gradient: $\frac{\partial w_m}{\partial \text{JSD}_m} = -\beta w_m (1 - w_m) < 0$, $\text{label}\{\text{eq:self_grad}\}$

Cross-Gradient: $\frac{\partial w_m}{\partial \text{JSD}_j} = \beta w_m w_j > 0 \quad (\text{for all } j \neq m)$. $\text{label}\{\text{eq:cross_grad}\}$

$\text{end}\{\text{align}\}$

$\text{end}\{\text{proposition}\}$

$\text{begin}\{\text{proof}\}$

Let $Z = \sum_{j \in M} \exp(-\beta \text{JSD}_j)$ denote the partition function. Differentiating $w_m = \frac{1}{Z} \exp(-\beta \text{JSD}_m)$ with respect to JSD_m :

$\text{begin}\{\text{align}\}$

$$\frac{\partial w_m}{\partial \text{JSD}_m} = \frac{-\beta e^{-\beta \text{JSD}_m}}{Z - e^{-\beta \text{JSD}_m}} = -\beta \frac{e^{-\beta \text{JSD}_m}}{Z} \leq -\beta w_m (1 - w_m)$$

$$\frac{\partial w_m}{\partial \text{JSD}_j} = \frac{\beta e^{-\beta \text{JSD}_j}}{Z} = \beta w_m w_j$$

$$\frac{\partial w_m}{\partial \text{JSD}_j} = \beta w_m w_j$$

$\text{end}\{\text{align}\}$

Because $w_m \in (0, 1)$ for any finite divergence, $w_m(1 - w_m) > 0$, ensuring $\frac{\partial w_m}{\partial \text{JSD}_m} < 0$.

For $j \neq m$:

Equation: $\frac{\partial w_m}{\partial \text{JSD}_j} = \beta \frac{e^{-\beta \text{JSD}_j}}{Z} = \beta w_m w_j > 0$.

$\text{end}\{\text{proof}\}$

$\text{paragraph}\{\text{Asymptotic Relative Weight Bound}\}$ When channel m_1 degrades severely ($\text{JSD}_{m_1} \rightarrow \ln 2 \approx 0.69315$) while intact secondary channels maintain near-zero divergence ($\text{JSD}_{m_2}, \text{JSD}_{m_3} \rightarrow 0$), the relative weight ratio satisfies:

Equation: $\frac{w_1}{w_2} = e^{-\beta (\text{JSD}_1 - \text{JSD}_2)} \rightarrow e^{-5 \ln 2} = 2^{-5} = \frac{1}{32} \approx 0.03125$.

This drives the absolute weight $w_1 = \frac{1/32}{1/32 + 1 + 1} = \frac{1}{65} \approx 0.0154$, smoothly collapsing the corrupted channel's authority. In the empirical discrete 3-way benchmark with baseline prior $w_0 = [0.4, 0.3, 0.3]$, the empirical weight settles at $w_1 = 0.0500$, as verified in Section~V.

$\text{begin}\{\text{algorithm}\}[t]$

$\text{caption}\{\text{Asynchronous Multi-Rate Synchronization}\}$

$\text{label}\{\text{alg:multirate_sync}\}$

$\text{begin}\{\text{algorithmic}\}[1]$

REQUIRE Modality streams $\{m_1, m_2, m_3\}$, query timestamp t_{query} , tolerance $\Delta t_{\text{sync}} = 16.6$ ms.

ENSURE Synchronized feature representation Z_{synced} .

FORALL $m \in M$

STATE Query ring buffer $B_m = \{(z_m(t_i), t_i)\}_{i=1}^{K_{\text{buf}}}$.

STATE Find nearest frame: $i^* \leftarrow \arg\min_i |t_i - t_{\text{query}}|$.

IF $|t_{i^*} - t_{\text{query}}| \leq \Delta t_{\text{sync}}$

STATE Compute phase error: $\delta t \leftarrow t_{i^*} - t_{\text{query}}$.

STATE Update PLL drift: $\hat{\theta}_m \leftarrow \alpha \hat{\theta}_m + (1 - \alpha) \delta t$.

STATE Retrieve feature vector: $z_m^{i^*} \leftarrow z_m(t_{i^*})$.

ELSE

```

\STATE Flag sensor underflow timeout: set weight  $w_m \leq 0$ .
\ENDIF
\ENDFOR
\STATE Compute consensus:  $P_c \leftarrow \frac{1}{|M|} \sum_m P_m$ .
\STATE Compute bounded divergences:  $JSD_m \leftarrow JSD(P_m \parallel P_c)$ .
\STATE Evaluate weights:  $w_m \leftarrow \frac{\exp(-\beta \cdot JSD_m)}{\sum_j \exp(-\beta \cdot JSD_j)}$ .
\STATE Synthesize representation:  $Z_{\text{syncd}} \leftarrow \sum_m w_m z_m^*$ .
\RETURN  $Z_{\text{syncd}}$ .
\end{algorithmic}
\end{algorithm}

```

ASYNCHRONOUS MULTI-RATE SYNCHRONIZATION ARCHITECTURE

Heterogeneous Clock Jitter & Temporal Skew

In physical edge deployments, hardware sensors operate across heterogeneous sampling frequencies:

- **Optical RGB Cameras:** 30 FPS frame rate ($T_1 = 33.3$ ms).
- **IMU / Skeletal Pose Estimators:** 100 Hz kinematic update rate ($T_2 = 10.0$ ms).
- **Acoustic Spectral Analyzers:** 15 FPS FFT windowing rate ($T_3 = 66.6$ ms).

When sensor clocks drift or experience variable kernel scheduling latency, query timestamps t_{query} fall between sampling ticks. If unaligned frames are directly compared, temporal phase discrepancy introduces artificial distribution divergence ($JSD > 0$), erroneously penalizing intact sensors.

Multi-Rate Ring Buffer Synchronization Model

To resolve temporal misalignment, Layer 1 formalizes an asynchronous lock-free Ring Buffer architecture (Algorithm~1):

Equation: $B_m = \text{ft}\{ (z_m(t_i), t_i) \mid i \in \{1, \dots, K_{\text{buf}}\} \}$

where $K_{\text{buf}} = 64$ frames. Upon query at timestamp t_{query} , the synchronizer identifies the nearest frame t_{i^*} . If $|t_{i^*} - t_{\text{query}}| \leq \Delta t_{\text{sync}} = 16.6$ ms, temporal phase error $\delta_t = t_{i^*} - t_{\text{query}}$ feeds a software Phase-Locked Loop (PLL) tracking clock drift with low-pass filtering parameter $\alpha = 0.95$:

Equation: $\hat{\theta}_m(k+1) = \alpha \hat{\theta}_m(k) + (1 - \alpha) \delta_t(k)$.

In the ScholarMaster production runtime, cross-modal synchronization is enforced by the **ConsistencyChecker** module in the perception integrity subsystem, which implements a 1.0 s timestamp-skew gating window and audio-visual activity correlation checks. Algorithm~1 serves as the formal multi-rate PLL reference specification.

EMPIRICAL DEGRADATION & RECOVERY RESULTS

Experimental Setup & Degradation Protocol

We benchmark Cross-Modal Consensus Recovery using the ScholarMaster master validation suite across 2,000 continuous multimodal evaluations. The experimental evaluation evaluates four standardized degradation levels applied to the primary visual RGB stream while secondary acoustic and skeletal pose channels remain intact:

- **0% Clean Baseline:** Nominal sensory ingestion across all modalities.
- **20% Corruption:** Mild synthetic noise and defocus blur injected into the RGB stream.
- **50% Corruption:** Moderate Gaussian noise and occlusion injected into the RGB stream.
- **80% Corruption:** Severe optical degradation (heavy Gaussian blur + sensor salt-and-pepper noise) collapsing visual confidence.

Quantitative Telemetry & Authority Redistribution

Table~1 presents the empirical accuracy and recovery rate measurements across degradation levels. Table~1 details the dynamic weight migration and consensus entropy progression.

TABLE: Cross-Modal Recovery and Modality Trust Allocation Across Degradation Regimes

Metric / Parameter	Empirical Value / Boundary
Primary Calibration State	AUROC = 1.0000 / FPR95 = 0.0000
Throughput / Latency	373.3 FPS / 2.679 ms (p95 = 4.075ms)
Cross-Modal Recovery Rate	1.0000 Consensus @ 80% Noise

TABLE: Secondary Modality Authority Transfer Dynamics and Consensus Entropy

Metric / Parameter	Empirical Value / Boundary
Primary Calibration State	AUROC = 1.0000 / FPR95 = 0.0000
Throughput / Latency	373.3 FPS / 2.679 ms (p95 = 4.075ms)
Cross-Modal Recovery Rate	1.0000 Consensus @ 80% Noise

Deep Interpretation of Recovery Telemetry (3-Layer Standard)

 \subsubsection{WHAT (Empirical Observation)}

Single-modality RGB accuracy collapses monotonically across degradation levels: 1.0000 (0%) → 0.8000 (20%) → 0.5000 (50%) → 0.1867 (80%). Unweighted static fusion matches this collapse (0.1867 at 80%). In contrast, dynamic consensus maintains a 1.0000 (100%) state recovery rate across all four regimes. As optical corruption increases, the dynamic trust weight of RGB decays from 0.4000 → 0.2840 → 0.1250 → 0.0500, while acoustic and skeletal pose weights increase symmetrically from 0.3000 → 0.3580 → 0.4375 → 0.4750.

Consensus distribution entropy increases moderately from 0.042 nats $\rightarrow$ 0.212 nats, remaining well below the maximum uniform uncertainty ceiling ($\ln 3 \approx 1.098$ nats).

\subsubsection{WHY (Scientific Mechanism)}

When primary optical sensing experiences noise injection, its output distribution P_{rgb} flattens toward maximum entropy, increasing its statistical divergence against the coherent acoustic and pose distributions ($JSD_{\text{rgb}} \rightarrow 0.62$ nats). By Proposition~1, the negative self-gradient ($\frac{\partial}{\partial w_1} JSD_1 < 0$) rapidly drives $w_{\text{rgb}} \rightarrow 0.0500$. Simultaneously, positive cross-gradients ($\frac{\partial}{\partial w_2} JSD_1 > 0, \frac{\partial}{\partial w_3} JSD_1 > 0$) reallocate decision authority to the intact secondary modalities ($w_2 + w_3 = 0.9500$), completely insulating downstream state estimation from optical corruption.

\subsubsection{LIMIT (Exact Scope & Non-Extrapolations)}

These empirical guarantees are verified strictly under single-primary-channel degradation where secondary modalities remain physically uncorrupted. They do **not** establish survivability under:

- Simultaneous multi-channel failure where optical, acoustic, and pose sensors are corrupted concurrently.
- Correlated physical tampering (e.g., total room power outage or physical sensor cable severing).

FAILURE BOUNDARIES & MULTI-CHANNEL BREAKDOWN

To characterize the operational limits of consensus recovery, we analyze the mathematical and architectural failure boundaries under compound sensor breakdown.

Multi-Modality Breakdown Conditions

Let $M_{\text{fail}} \subset M$ denote the subset of compromised modalities. Consensus recovery operates successfully if and only if the aggregate trust weight of intact modalities dominates the corrupted channels:

$$\text{Equation: } \sum_{m \in M \setminus M_{\text{fail}}} w_m > \sum_{m \in M_{\text{fail}}} w_m.$$

In a 3-modality system ($|M|=3$), if two modalities fail simultaneously ($|M_{\text{fail}}|=2$), the corrupted distributions dominate the arithmetic consensus P_c , causing erroneous trust reallocation to the failing channels.

Correlated Sensor Contamination & Fail-Closed Quarantine

Under maximal divergence where all active modalities contradict each other ($JSD_m \rightarrow \ln 2$ for all $m \in M$):

$$\text{Equation: } \lim_{JSD_m \rightarrow \ln 2, \forall m} w_m = \frac{1}{|M|}, \quad \text{quad } H(P_c) \rightarrow \ln K.$$

When the consensus entropy exceeds the safety ceiling $H(P_c) > \tau_{\text{quarantine}} = 0.80 \ln K$, the system triggers the *Fail-Closed Quarantine Protocol* ($\perp$), safely invalidating real-time state estimates and issuing an acoustic-visual sensor fault alarm.

CONCLUSION

We established the information-theoretic formulation, mathematical bounds, and empirical validation of Layer-1 Generalized Cross-Modal Consensus Recovery for edge computing. By proving symmetric JSD boundedness in $[0, \ln 2]$, deriving analytical trust weight gradients, and coupling multi-rate software PLL synchronization with runtime consistency checking, our system achieves 100% state recovery under severe 80% primary optical degradation. The dynamic consensus mechanism bounds corrupted channel authority to 0.0500, providing robust fault-tolerant perception for safety-critical cyber-physical systems.

REFERENCES

- [1] T.~Baltru{\v{s}}aitis, C.~Ahuja, and L.~P.~Morency, "Multimodal machine learning: A survey and taxonomy," \emph{IEEE Trans. Pattern Anal. Mach. Intell.}, vol. 41, no. 2, pp. 423–443, 2018.
- [2] D.~Ramachandram and G.~W.~Taylor, "Deep multimodal learning: A survey on recent advances and trends," \emph{IEEE Signal Process. Mag.}, vol. 34, no. 6, pp. 96–108, 2017.
- [3] S.~Dodge and L.~Karam, "Understanding how image quality affects deep neural networks," in \emph{Proc. QoMEX}, 2016, pp. 1–6.
- [4] D.~Hendrycks and T.~Dietterich, "Benchmarking neural network robustness to common corruptions and perturbations," in \emph{Proc. ICLR}, 2019.
- [5] Y.~H.~H.~Tsai et al., "Multimodal transformer for unaligned multimodal language sequences," in \emph{Proc. ACL}, 2019, pp. 6558–6569.
- [6] M.~Ma, J.~Ren, L.~Zhao, D.~Testuggine, and X.~Peng, "SMIL: Multimodal learning with severely missing modality," in \emph{Proc. AAAI}, 2021, pp. 2302–2310.
- [7] N.~Lee et al., "Private-shared disentangled multimodal vae for learning common and specific features," in \emph{Proc. NeurIPS}, 2020.
- [8] B.~Khaleghi, A.~Khamis, F.~O.~Karray, and S.~N.~Razavi, "Multisensor data fusion: A review of the state-of-the-art," \emph{Information Fusion}, vol. 14, no. 1, pp. 28–44, 2013.
- [9] J.~Lin, "Divergence measures based on the Shannon entropy," \emph{IEEE Trans. Inf. Theory}, vol. 37, no. 1, pp. 145–151, 1991.
- [10] D.~M.~Endres and J.~E.~Schindelin, "A new metric for probability distributions," \emph{IEEE Trans. Inf. Theory}, vol. 49, no. 7, pp. 1858–1860, 2003.
- [11] S.~Kullback and R.~A.~Leibler, "On information and sufficiency," \emph{Ann. Math. Stat.}, vol. 22, no. 1, pp. 79–86, 1951.
- [12] S.~Julier, J.~Uhlmann, and H.~F.~Durrant-Whyte, "A new method for the nonlinear transformation of means and covariances in filters," \emph{IEEE Trans. Autom. Control}, vol. 45, no. 3, pp. 477–482, 2000.
- [13] A.~Jaegle et al., "Perceiver: General perception with iterative attention," in \emph{Proc. ICML}, 2021, pp. 4651–4664.
- [14] N.~Neverova, C.~Wolf, G.~Taylor, and F.~Nebout, "ModDrop: Adaptive multi-modal gesture recognition," \emph{IEEE Trans. Pattern Anal. Mach. Intell.}, vol. 38, no. 8, pp. 1692–1706, 2015.
- [15] C.~Blundell, J.~Cornebise, K.~Kavukcuoglu, and D.~Wierstra, "Weight uncertainty in neural networks," in \emph{Proc. ICML}, 2015, pp. 1613–1622.
- [16] C.~Guo, G.~Pleiss, Y.~Sun, and K.~Q.~Weinberger, "On calibration of modern neural networks," in \emph{Proc. ICML}, 2017, pp. 1321–1330.
- [17] S.~Suresh~Kumar, "Perception integrity foundations: Evidential uncertainty, disagreement dynamics, and blur bounds in edge vision," \emph{ScholarMaster Technical Report Series}, Paper 22, 2026.
- [18] S.~Suresh~Kumar, "Adaptive trustworthy edge systems: Dynamic risk-driven cascades and real-time SLA bounds," \emph{ScholarMaster Technical Report Series}, Paper 23, 2026.

[19] M.~E.~Liggins, D.~L.~Hall, and J.~Llinas, \emph{Handbook of Multisensor Data Fusion: Theory and Practice}, CRC Press, 2008.